



KubeCon



CloudNativeCon

Europe 2024



Why Kubernetes Is Inappropriate for Platforms and How to Make It Better.

 Stefan Schimanski, Upbound, Senior Principal Engineer

 Mangirdas Judeikis, CAST AI, Staff Engineer

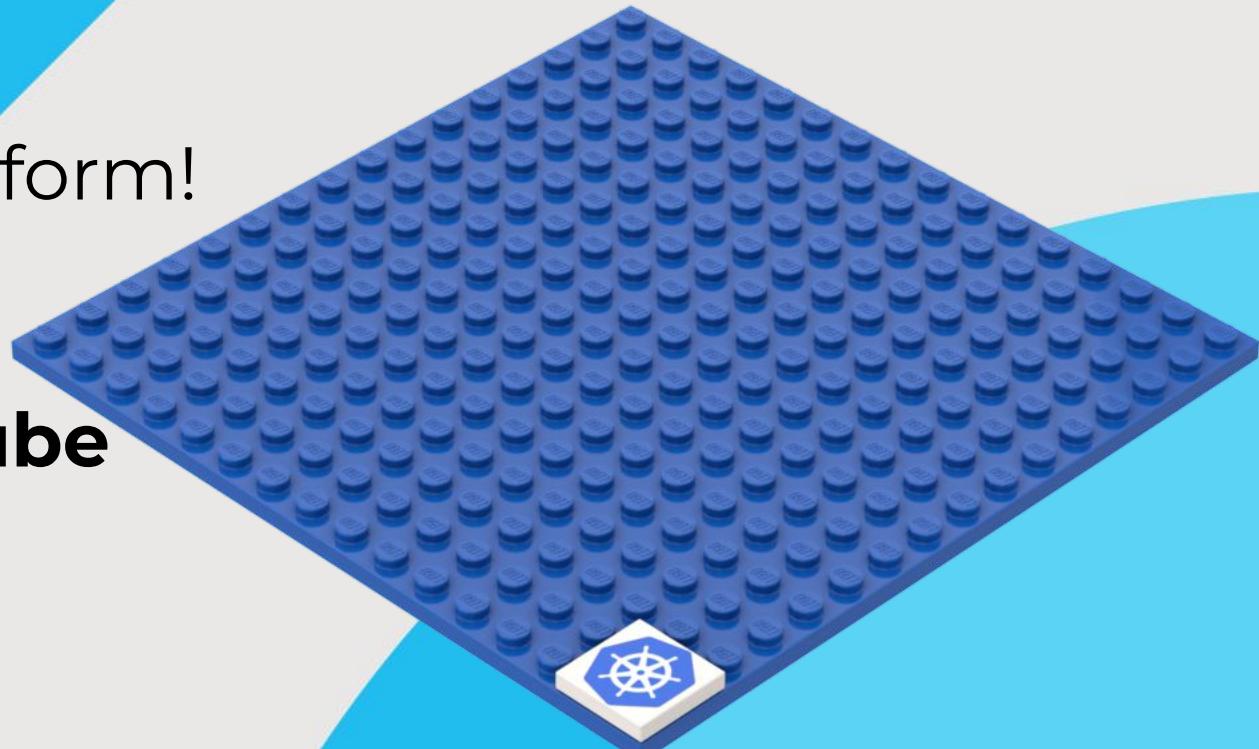
 Sebastian Scheele, Kubermatic - CEO



Thought experiment for today

Let's build an internal AI/ML platform!

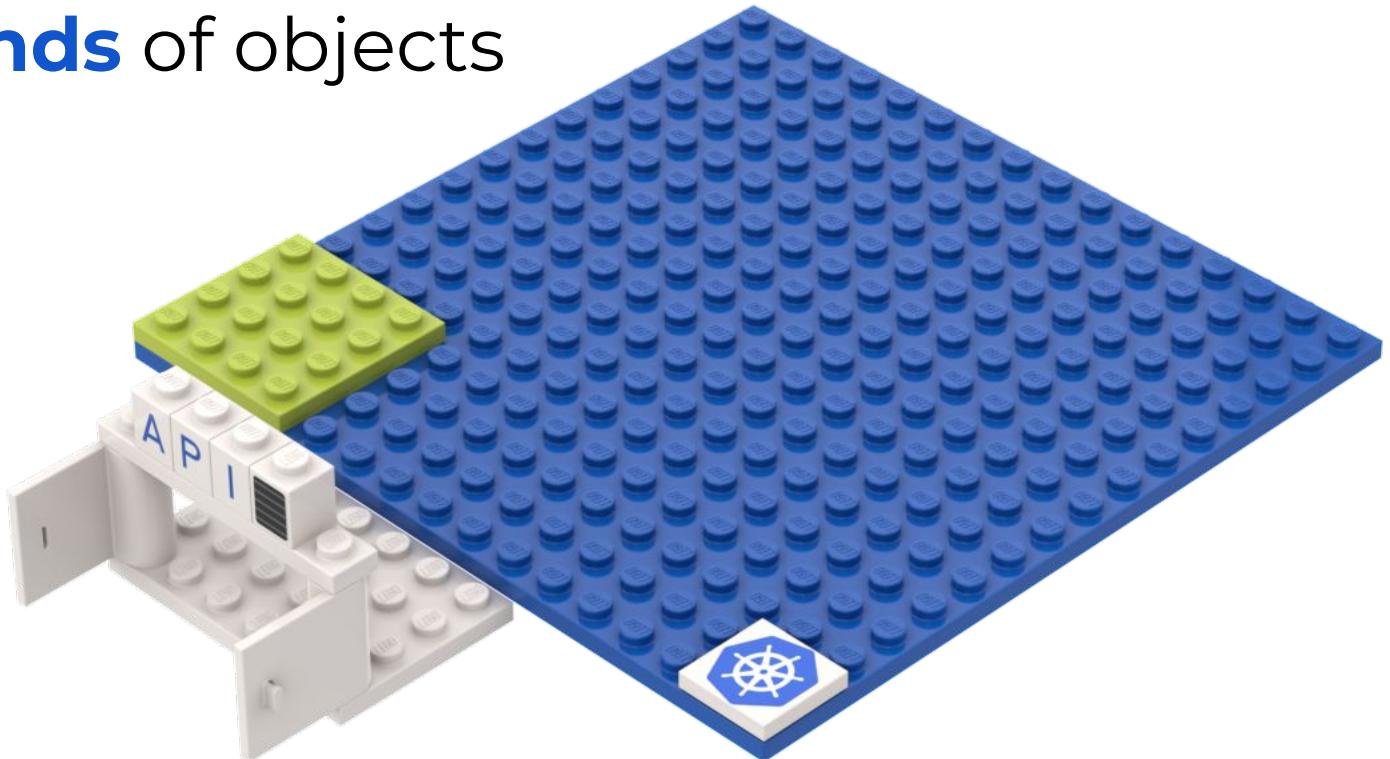
... and remind us **why we like kube**
... and **what we dislike.**



1. An Object Universe

```
kind: Model
apiVersion: training.faros.sh/v1alpha1
metadata:
  name: my-amazing-chat
  namespace: team1
spec:
  model: Llama2
  nProcPerNod: 1
  script: chat_completion.py
  ckptDir: /pvc/data
  tokenizerPath: tokenizer.model
  maxSeqLen: 512
  maxBatchSize: 6
```

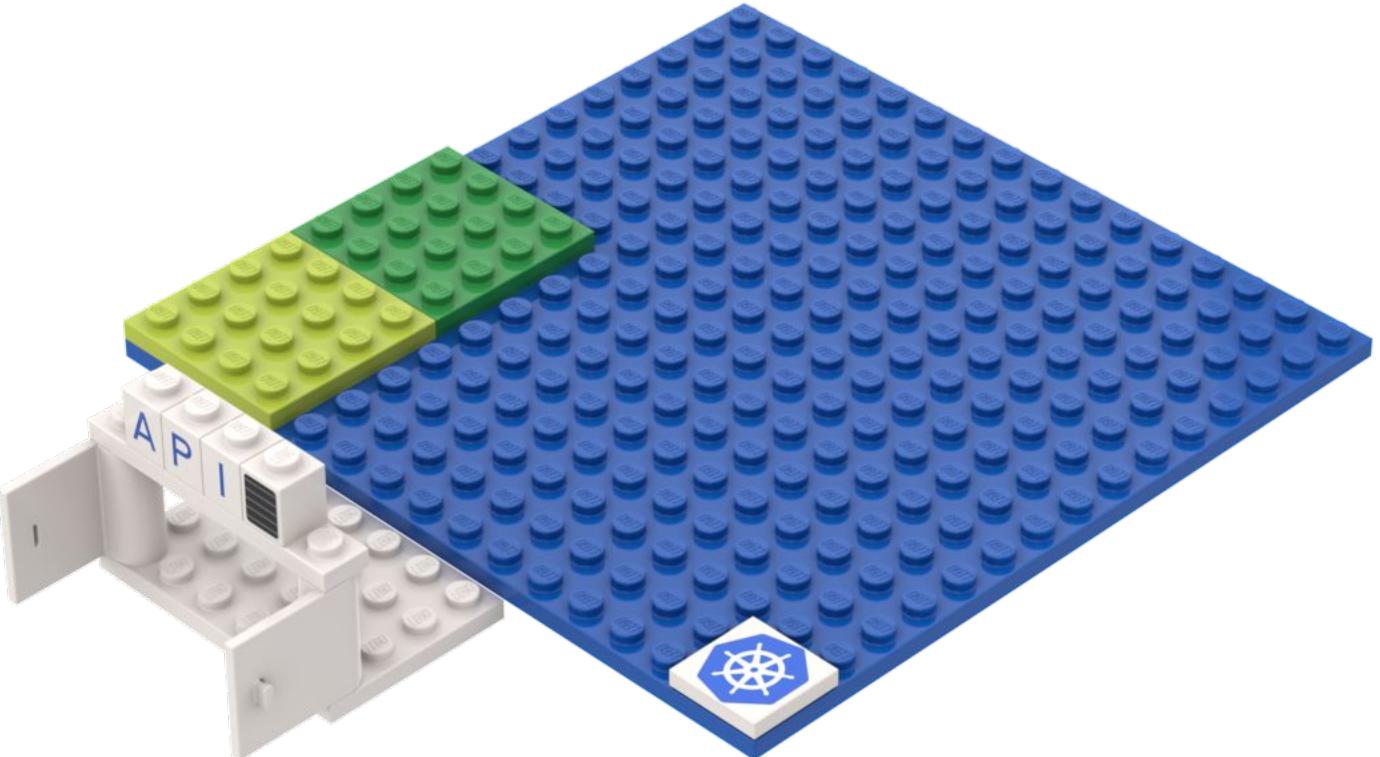
Different **kinds** of objects



1. An Object Universe

```
kind: Model
apiVersion: training.faros.sh/v1alpha1
metadata:
  name: my-amazing-chat
  namespace: team1
spec:
  model: Llama2
  nProcPerNod: 1
  script: chat_completion.py
  ckptDir: /pvc/data
  tokenizerPath: tokenizer.model
  maxSeqLen: 512
  maxBatchSize: 6
```

```
kind: Model
apiVersion: training.faros.sh/v1alpha1
metadata:
  name: my-amazing-chat
  namespace: team2
spec:
  model: Llama2
  nProcPerNod: 2
  script: chat_completion.py
  ckptDir: /pvc/data
  tokenizerPath: tokenizer.model
  maxSeqLen: 512
  maxBatchSize: 6
```



Living in **namespaces**.

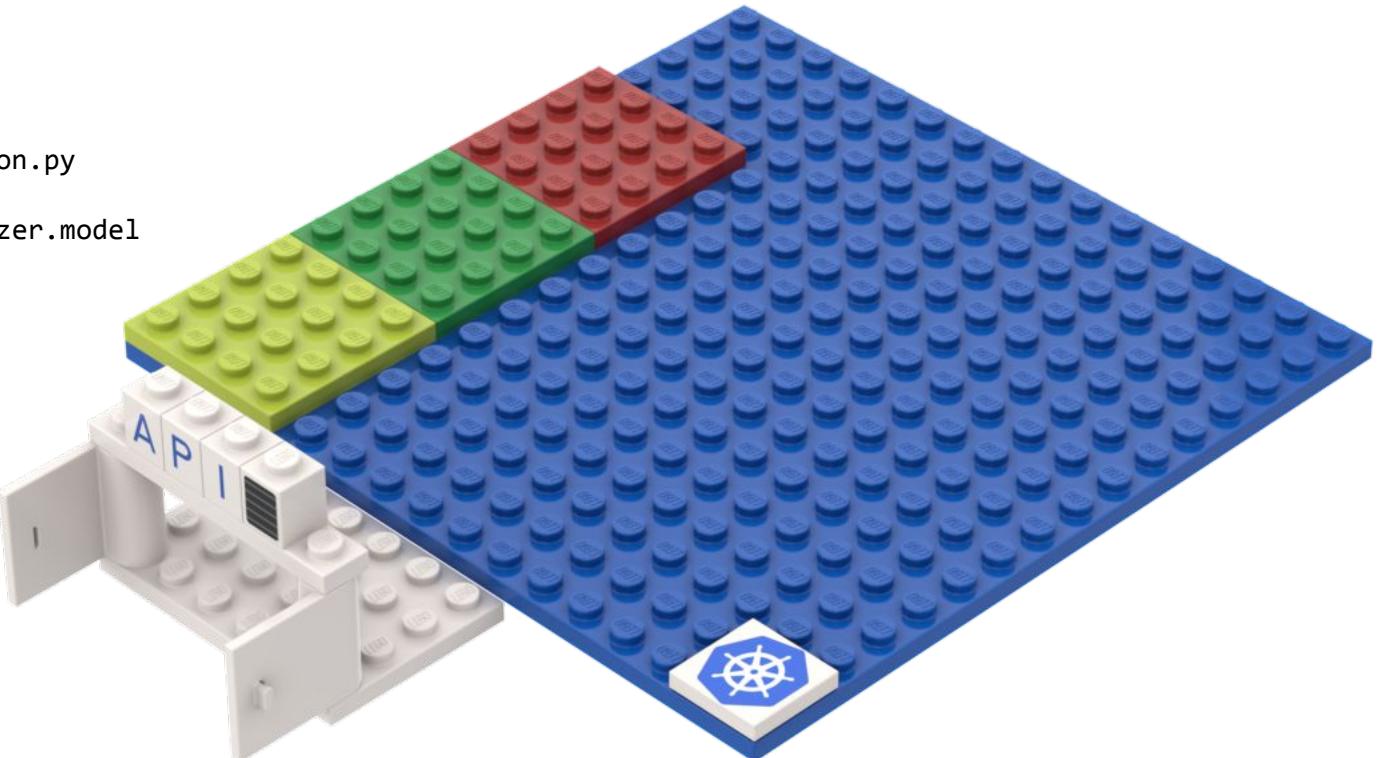
1. An Object Universe

```
kind: Model
apiVersion: training.faros.sh/v1alpha1
metadata:
  name: my-amazing-chat
  namespace: team1
spec:
  model: Llama2
  nProcPerNod: 1
  script: chat_completion.py
  ckptDir: /pvc/data
  tokenizerPath: tokenizer.model
  maxSeqLen: 512
  maxBatchSize: 6
```

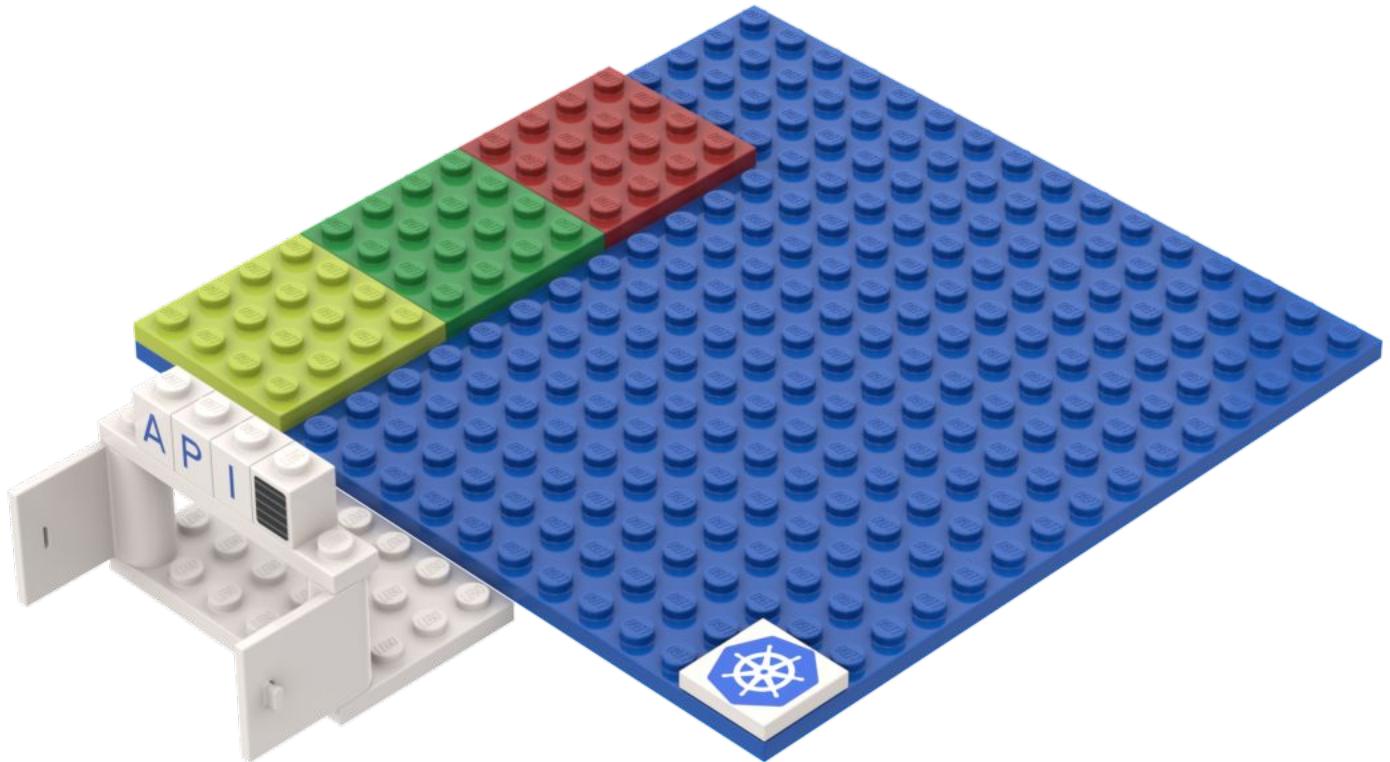
```
kind: Model
apiVersion: training.faros.sh/v1alpha1
metadata:
  name: my-amazing-chat
  namespace: team2
spec:
  model: Llama2
  nProcPerNod: 2
  script: chat_completion.py
  ckptDir: /pvc/data
  tokenizerPath: tokenizer.model
  maxSeqLen: 512
  maxBatchSize: 6
```

```
kind: Model
apiVersion: training.faros.sh/v1alpha1
metadata:
  name: my-amazing-chat-internal
  namespace: team3
spec:
  model: Llama2
  nProcPerNod: 2
  script: chat_completion.py
  ckptDir: /pvc/data
  tokenizerPath: tokenizer.model
  maxSeqLen: 512
  maxBatchSize: 3
```

with a **uniform API**.

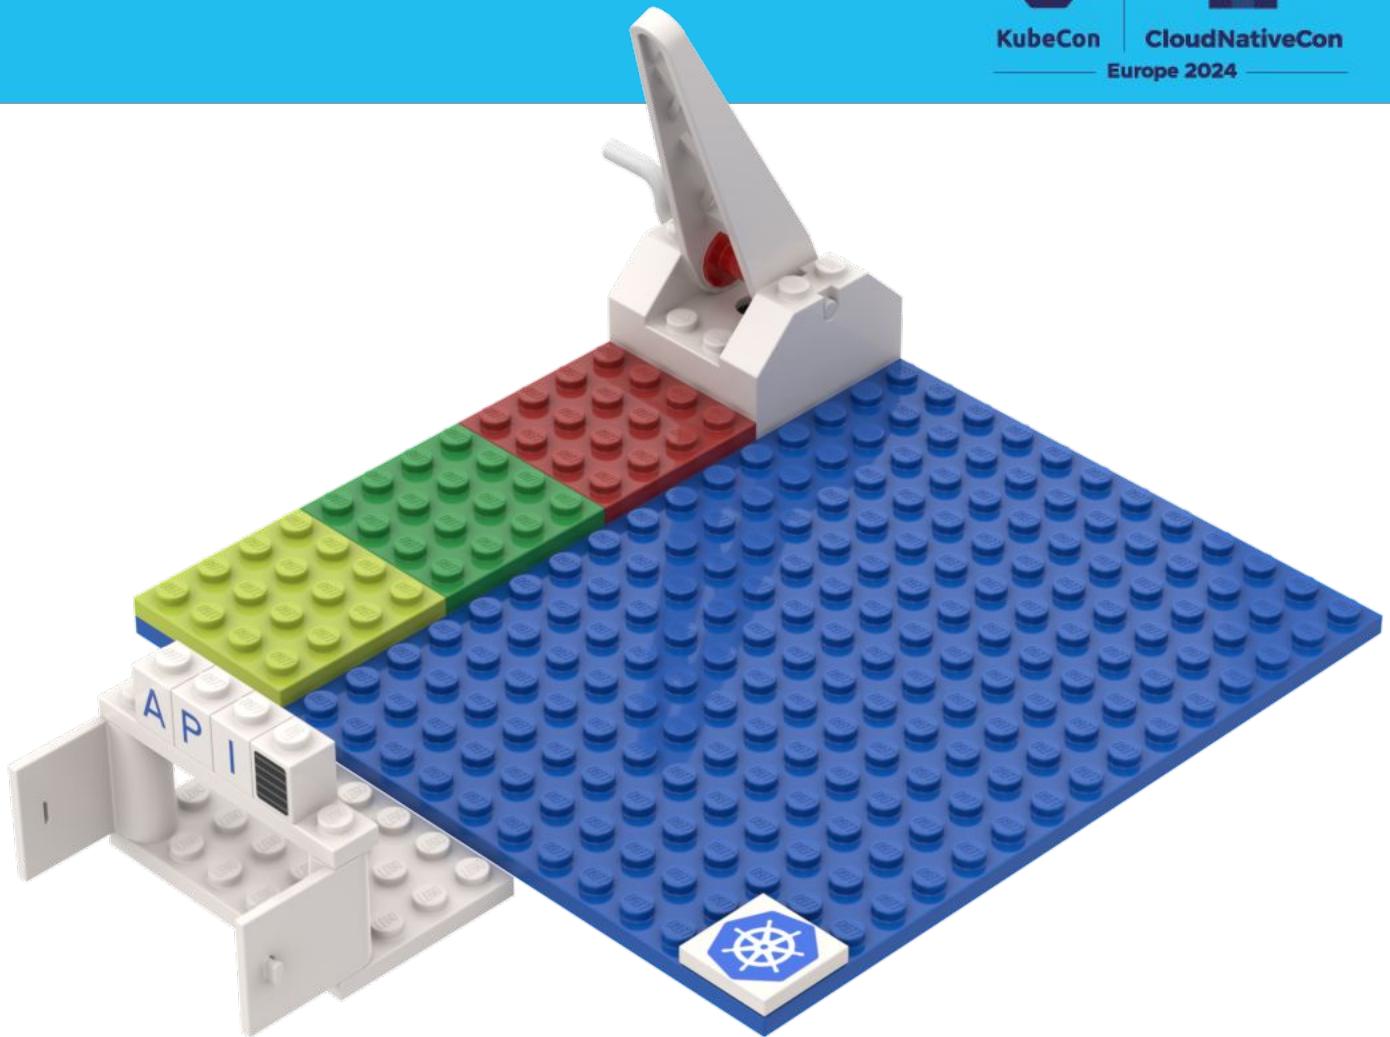


2. The Reconciler Pattern



2. The Reconciler Pattern

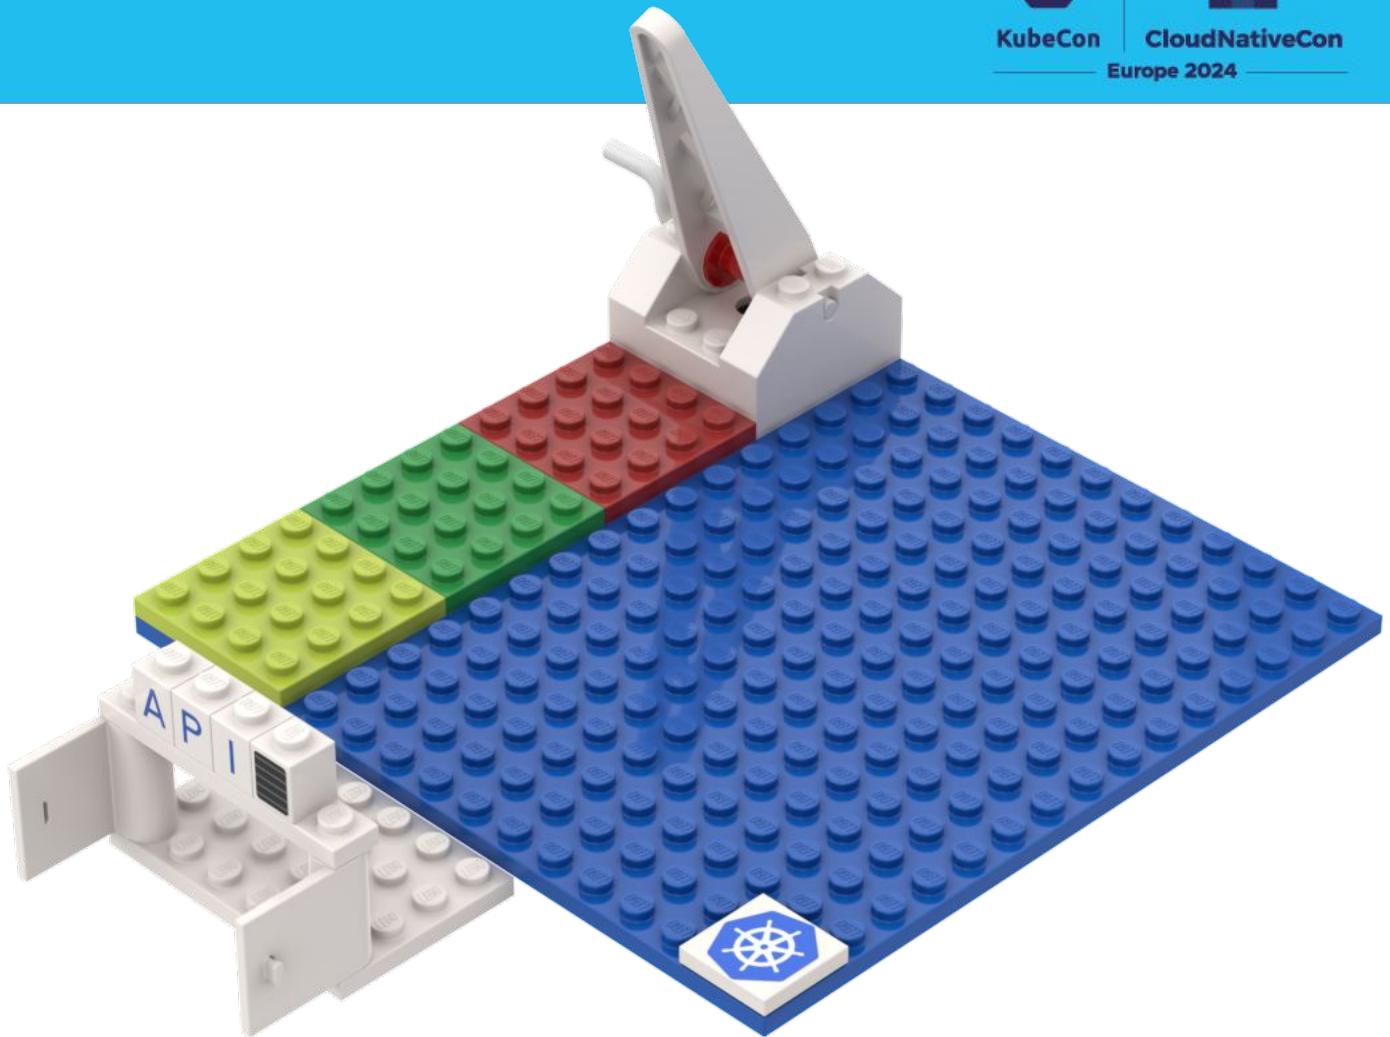
Controller runtime



2. The Reconciler Pattern

Controller runtime

kubebuilder



2. The Reconciler Pattern

Controller runtime

kubebuilder

metacontroller



2. The Reconciler Pattern

Controller runtime

kubebuilder

metacontroller

client-go



2. The Reconciler Pattern

Controller runtime

kubebuilder

metacontroller

client-go

kube-rs, python, java, ...



Big ecosystem of tooling

3. Multi-Tenancy?



3. Multi-Tenancy?

Namespaces

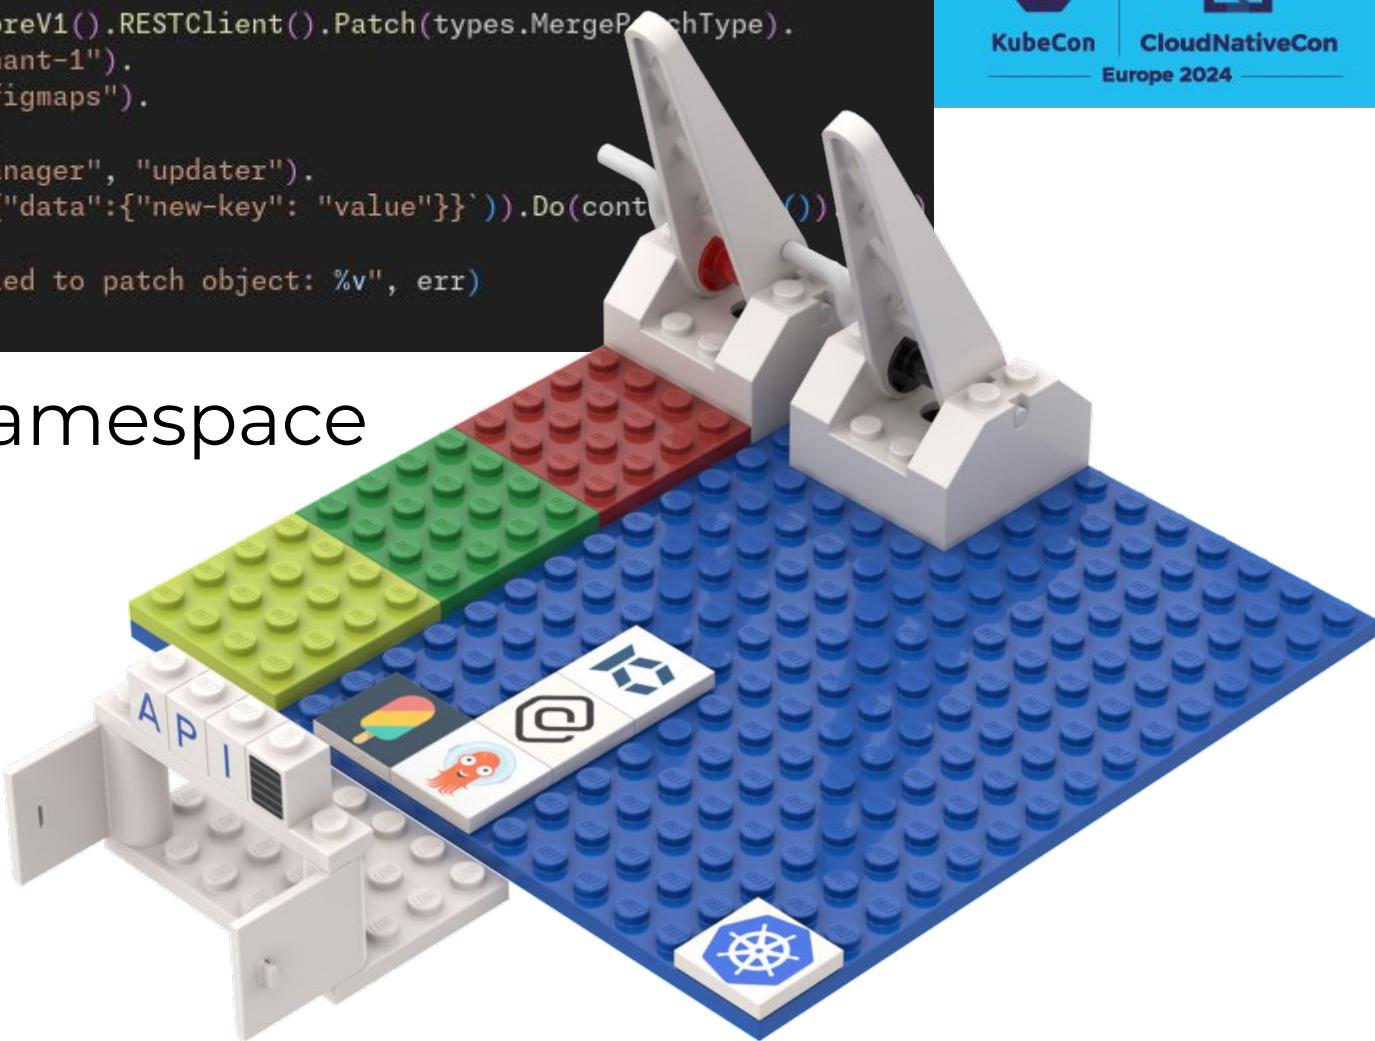


3. Multi-Tenancy?

Namespaces

```
_, err = client.CoreV1().RESTClient().Patch(types.MergePatchType).  
    Namespace("tenant-1").  
    Resource("configmaps").  
    Name("my-cm").  
    Param("fieldManager", "updater").  
    Body([]byte(`{"data":{"new-key": "value"}`)).Do(context.Background())  
if err != nil {  
    t.Fatalf("Failed to patch object: %v", err)  
}
```

Clients access all or one namespace



3. Multi-Tenancy?

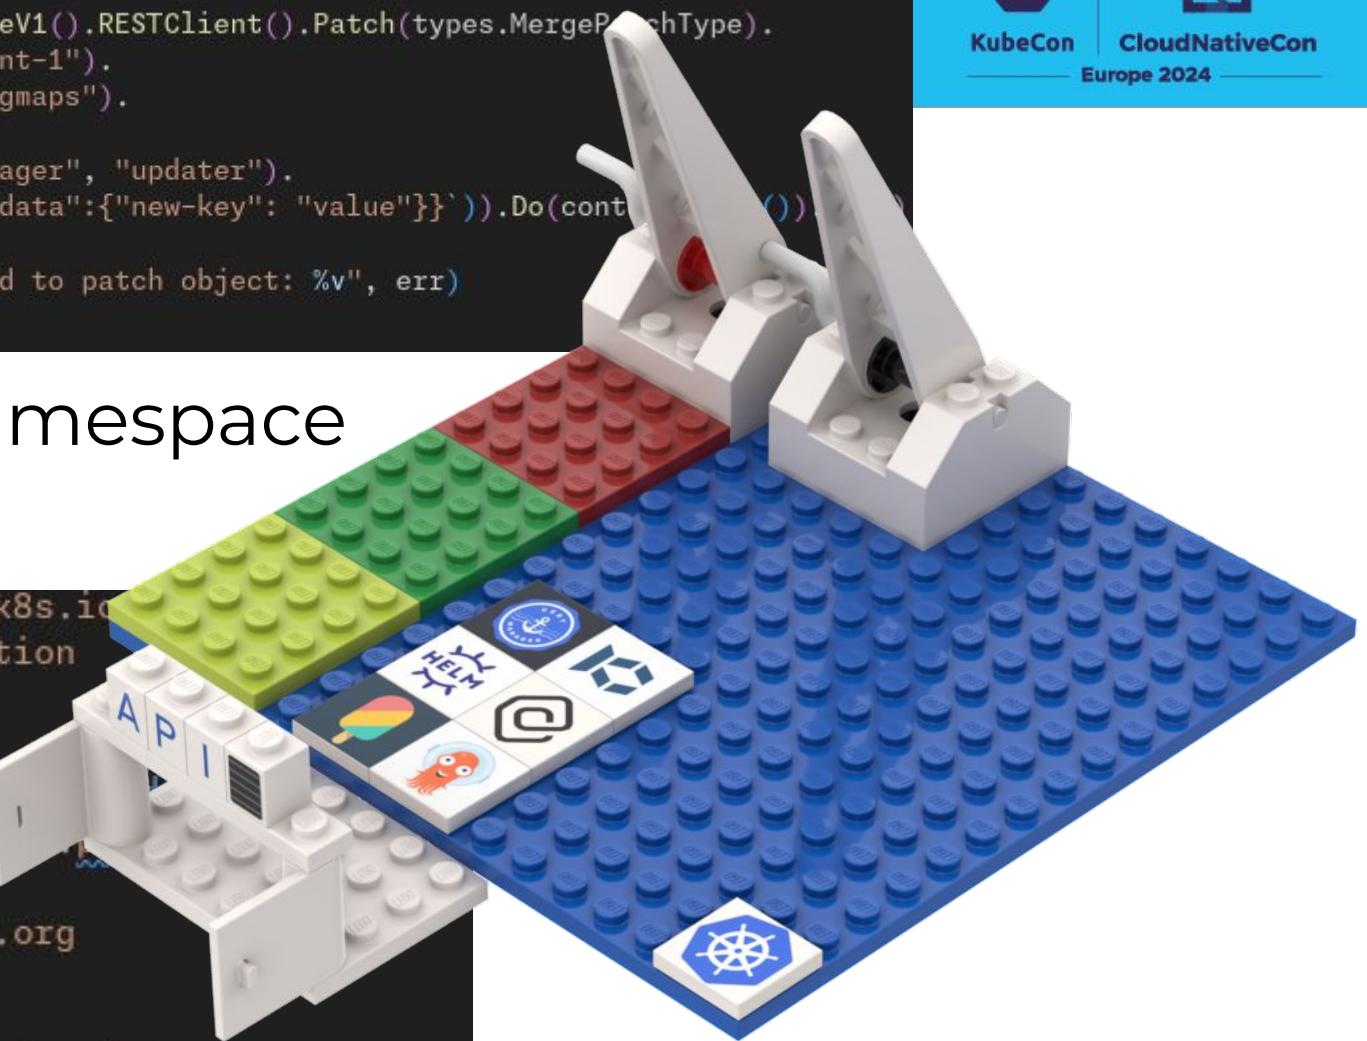
Namespaces

```
_ , err = client.CoreV1().RESTClient().Patch(types.MergePatchType).  
Namespace("tenant-1").  
Resource("configmaps").  
Name("my-cm").  
Param("fieldManager", "updater").  
Body([]byte(`{"data":{"new-key": "value"}}`)).Do(context.Background())  
if err != nil {  
    t.Fatalf("Failed to patch object: %v", err)  
}
```

Clients access all or one namespace

CRDs

```
apiVersion: apiextensions.k8s.io/v1  
kind: CustomResourceDefinition  
metadata:  
  labels:  
    addonmanager.kubernetes.io/mode: Reconcile  
    name: bgpconfigurations.  
spec:  
  group: crd.projectcalico.org  
  names:  
    kind: BGPConfiguration  
    listKind: BGPConfigurationList  
    plural: bgpconfigurations  
    singular: bgpconfiguration  
  scope: Cluster  
  versions:  
  - name: v1
```



3. Multi-Tenancy?



KubeCon

CloudNativeCon
Europe 2024

Namespaces

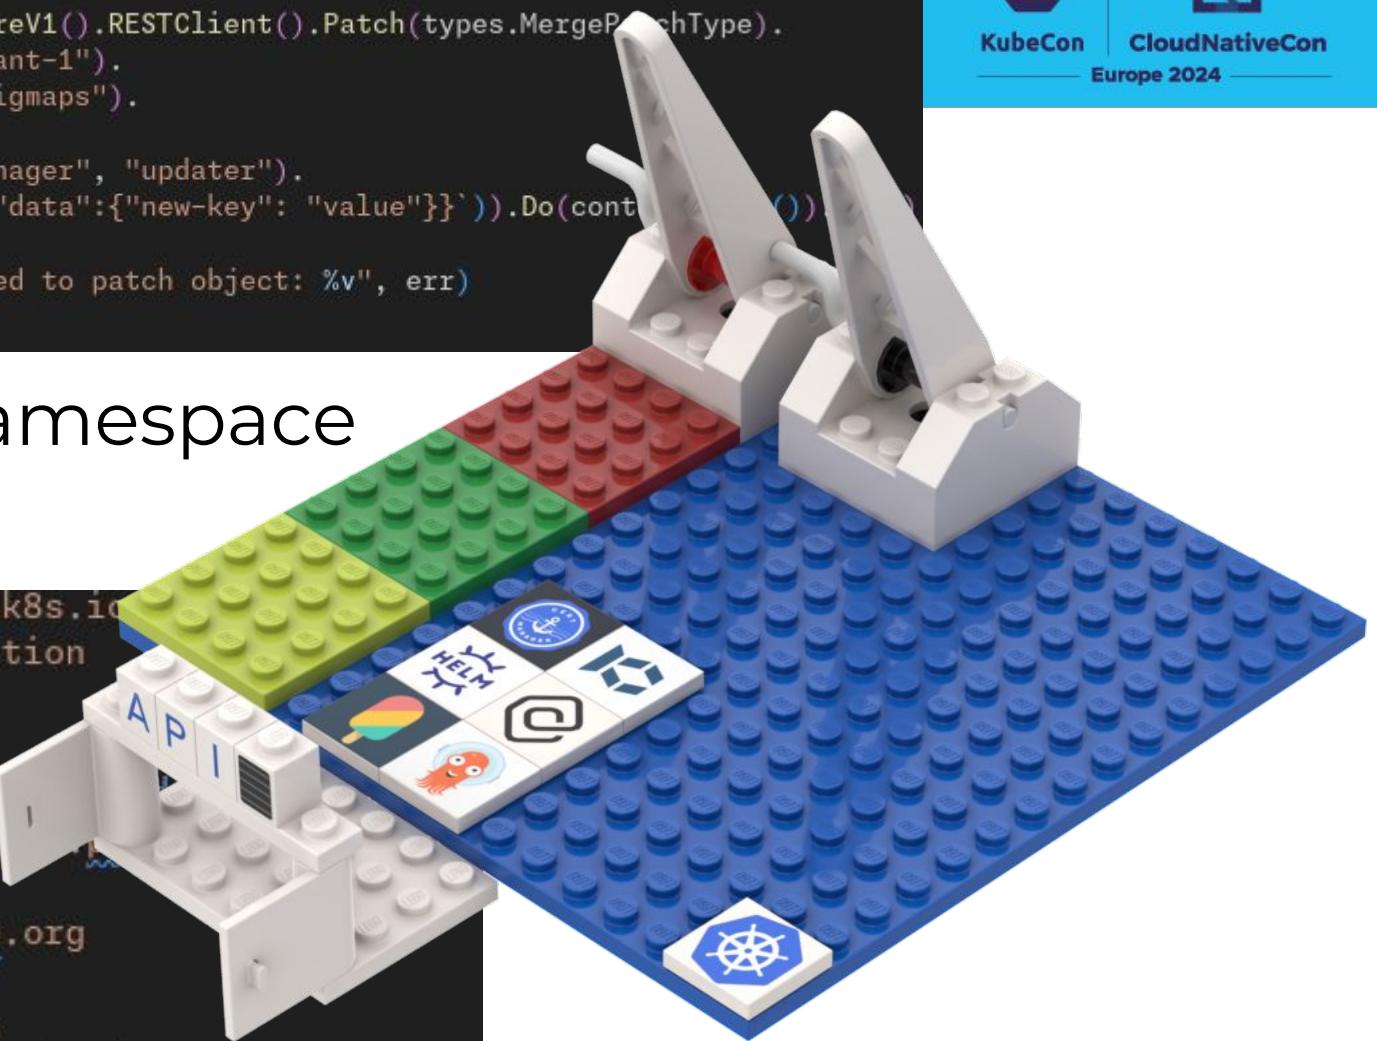
```
_ , err = client.CoreV1().RESTClient().Patch(types.MergePatchType).  
Namespace("tenant-1").  
Resource("configmaps").  
Name("my-cm").  
Param("fieldManager", "updater").  
Body([]byte(`{"data":{"new-key": "value"}`))).Do(context.Background())  
if err != nil {  
    t.Fatalf("Failed to patch object: %v", err)  
}
```

Clients access all or one namespace

CRDs



```
apiVersion: apiextensions.k8s.io/v1  
kind: CustomResourceDefinition  
metadata:  
  labels:  
    addonmanager.kubernetes.io/mode: Reconcile  
    name: bgpconfigurations.  
spec:  
  group: crd.projectcalico.org  
  names:  
    kind: BGPConfiguration  
    listKind: BGPConfigurationList  
    plural: bgpconfigurations  
    singular: bgpconfiguration  
  scope: Cluster  
  versions:  
    - name: v1
```



Short Status Check

This looks ok'ish.
Cracks already visible.

But we have learned to
live with our pains.

Let's continue building!



This picture is very 2017

Single region

Multi-tenancy via namespaces

Single source of truth



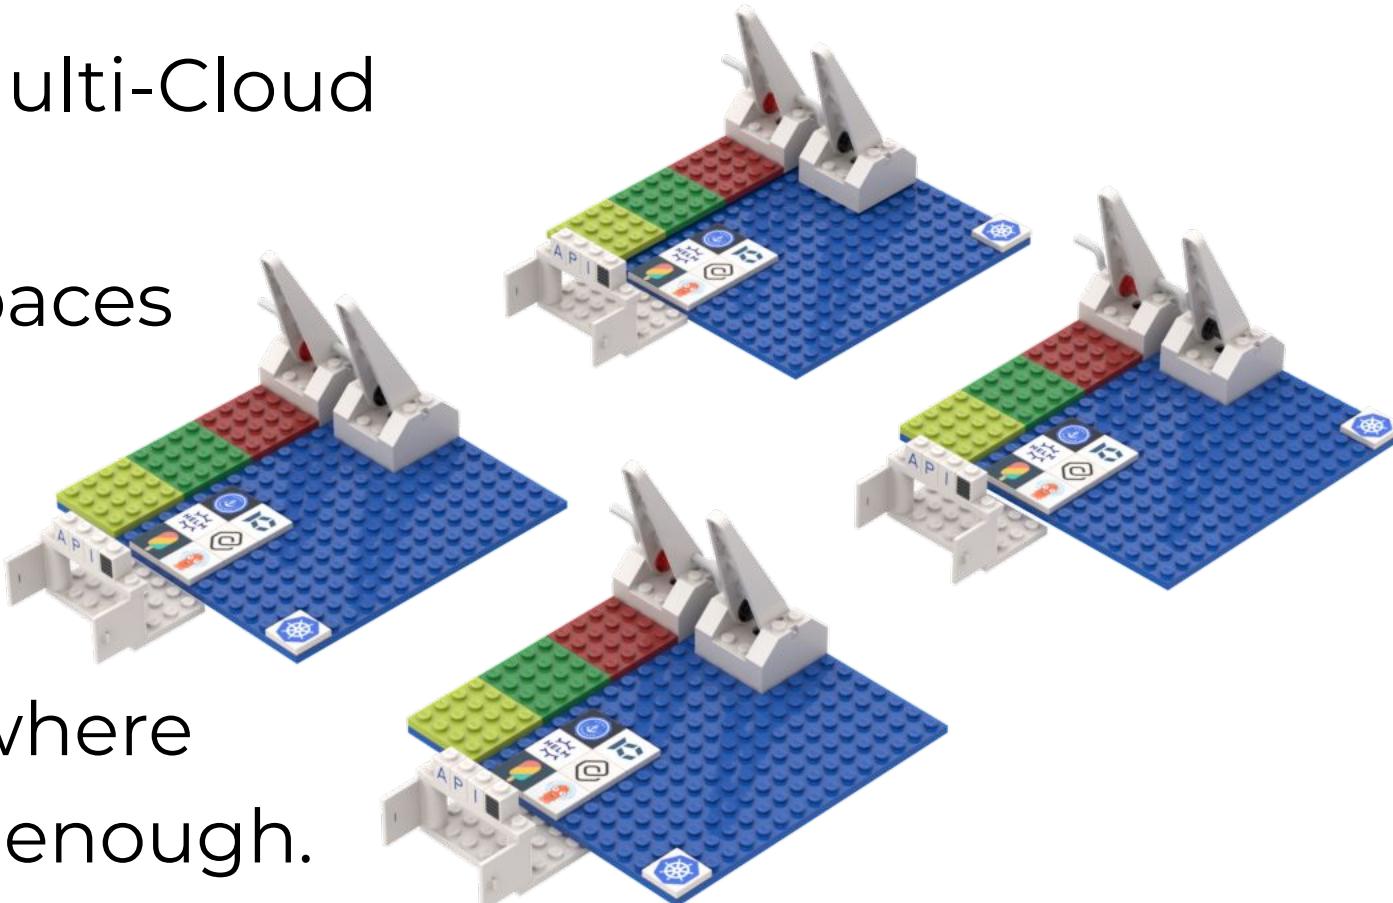
~~Single region~~ Multi-region Multi-Cloud

Multi-tenancy via namespaces

~~Single source of truth~~

Even more clusters where namespaces are not enough.

Sprawl of clusters.



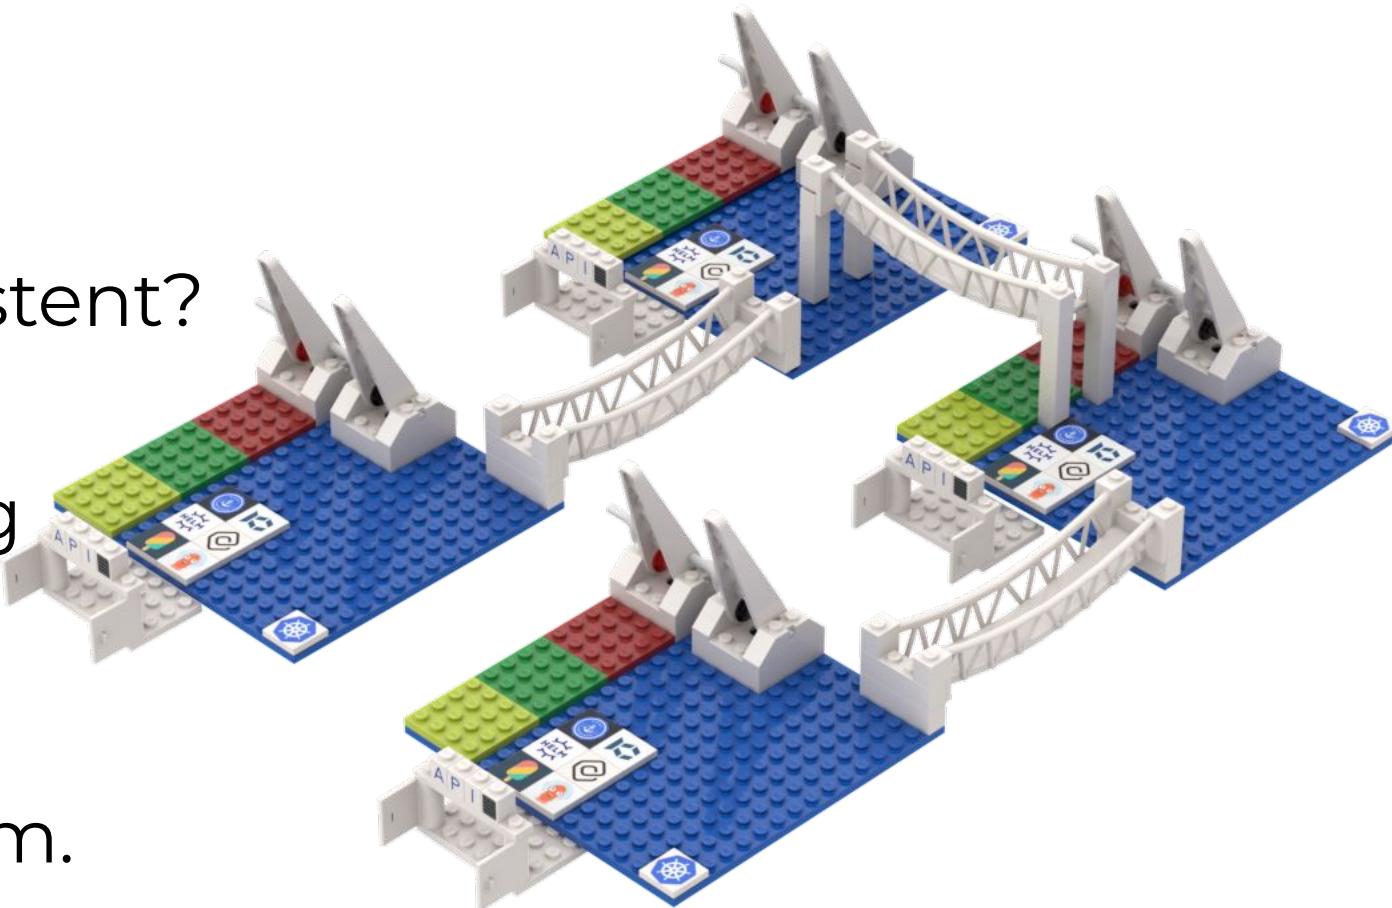
... with that more challenges.

How to share data?

How to keep config consistent?

How to keep everything
compliant?

We need bridges, lots of them.



We have multi-cluster tools – many to create them

vclusters

Kamaji

KubeVirt

Karmada

Cluster API



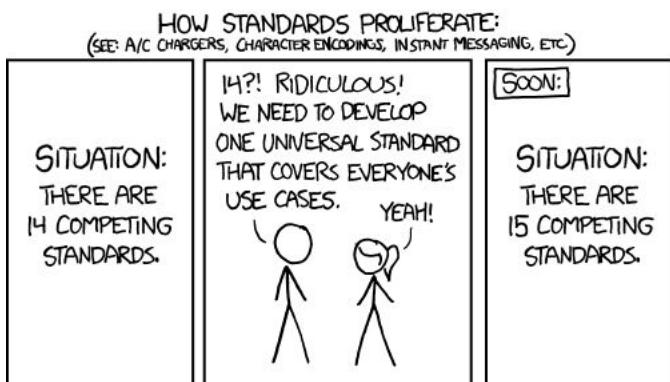
We have multi-cluster tools – many to tame them

ArgoCD

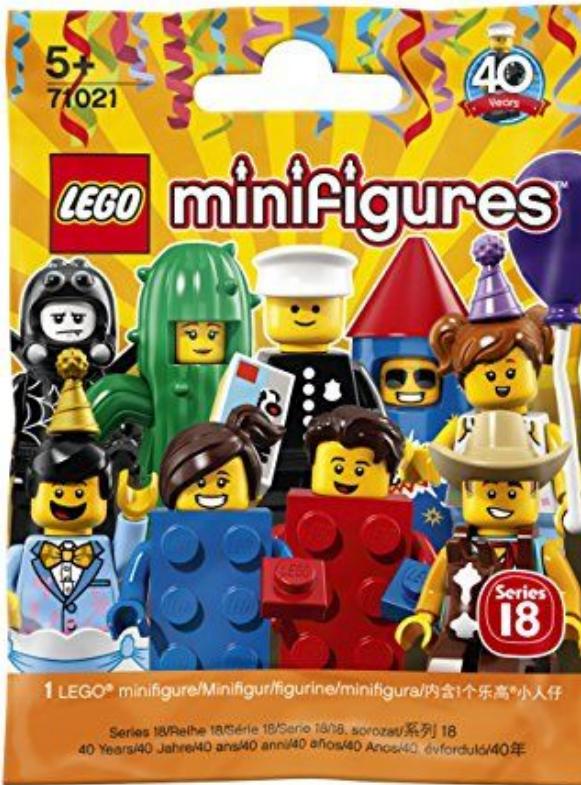
Crossplane

Cluster Federation

OpenClusterManager



But technology and tools are not the only problem.



In platforms, we have different personas.

And every single tool dictates their view on these.

Roles & responsibilities - permissions



Roles & responsibilities - permissions

Typical Roles & responsibilities:

- platform owner
- service provider
- user

Everyone with partial responsibility.

And finally 3rd-parties.

Everybody can build a helm chart for one cluster.

How to package something for this setup?



Complexity explodes

hard to support

hard to extend

hard to integrate into

not a good experience
for any of the roles.



Reality Check

Kube was built to run containers.

Kube was never built for platforms.

Reality Check

Kube was built to run containers.

Kube was never built for platforms.

Remember why we like Kube.

Reality Check

Kube was built to run containers.

Kube was never built for platforms.

Remember why we like Kube.

What would be a Kube built for platforms?

Towards a Kubernetes for Platforms

We know our final target. We did it before. But this time:

Ambitions are different.
Personas changed.
Kube was never the endgame.

Kubernetes is a platform for building platforms. It's a better place to start; not the endgame.

22:04 · 27.11.17 · Twitter Web Client

248 Retweets 39 Zitierte Tweets 735 „Gefällt mir“-Angaben

A necessary shift in mindset



Traditional

lots of fun, great for creativity, but everything built out of these will look differently. Generally too low level.



A more Grown up Approach

Everything lives in one environment. We want off-the-shelf components. We want one environment where we can collaborate. Where solution from one party fit to setup of somebody else.

Three Personas – Three Views



The platform owner
the ones with the keys
for everything,
the ones connecting and aligning
everything



The service provider
AI/ML service team,
DB team, PKI team,
CI/CD service, 3rd-parties
service owner



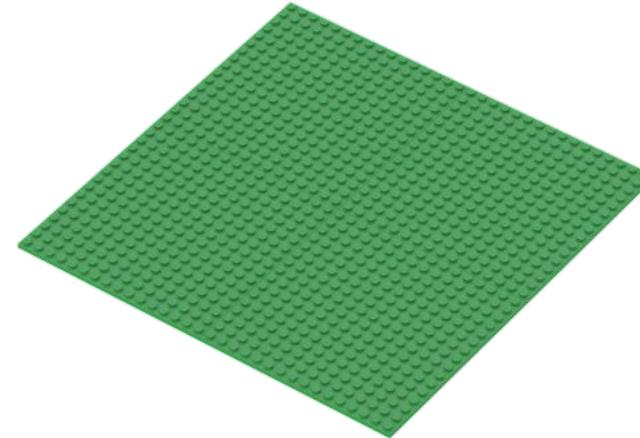
The users
Developer,
Data scientist,
application owner



The platform owner
the ones with the keys
for everything,
the ones connecting and aligning
everything

A Kubernetes for Platforms – the Platform Owner

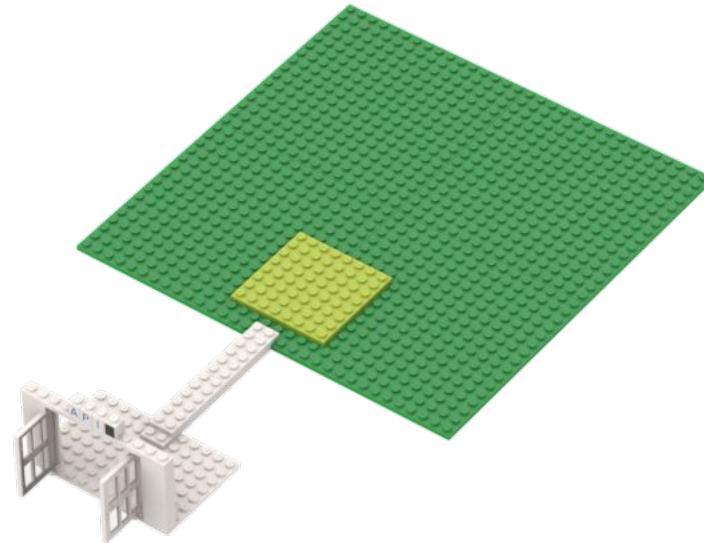
As a platform owner - **you are an enabler!**



A Kubernetes for Platforms – the Platform Owner

As a platform owner - you are an enabler!

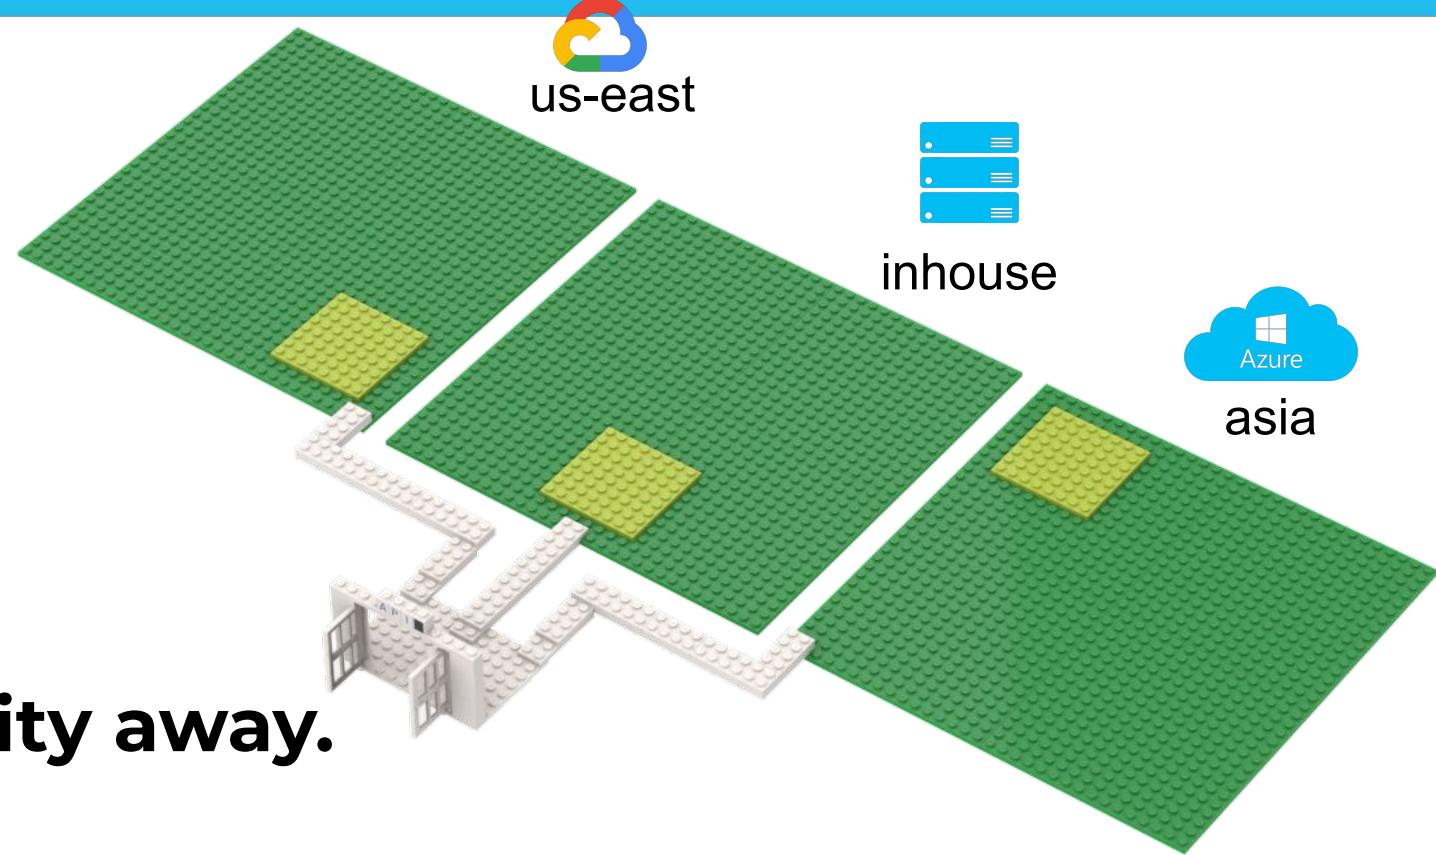
Give user a **well defined, flexible platform to ease their job!**



A Kubernetes for Platforms – the Platform Owner

As a platform owner

Give user a well defined,
flexible platform
to ease their job!



Abstract the complexity away.



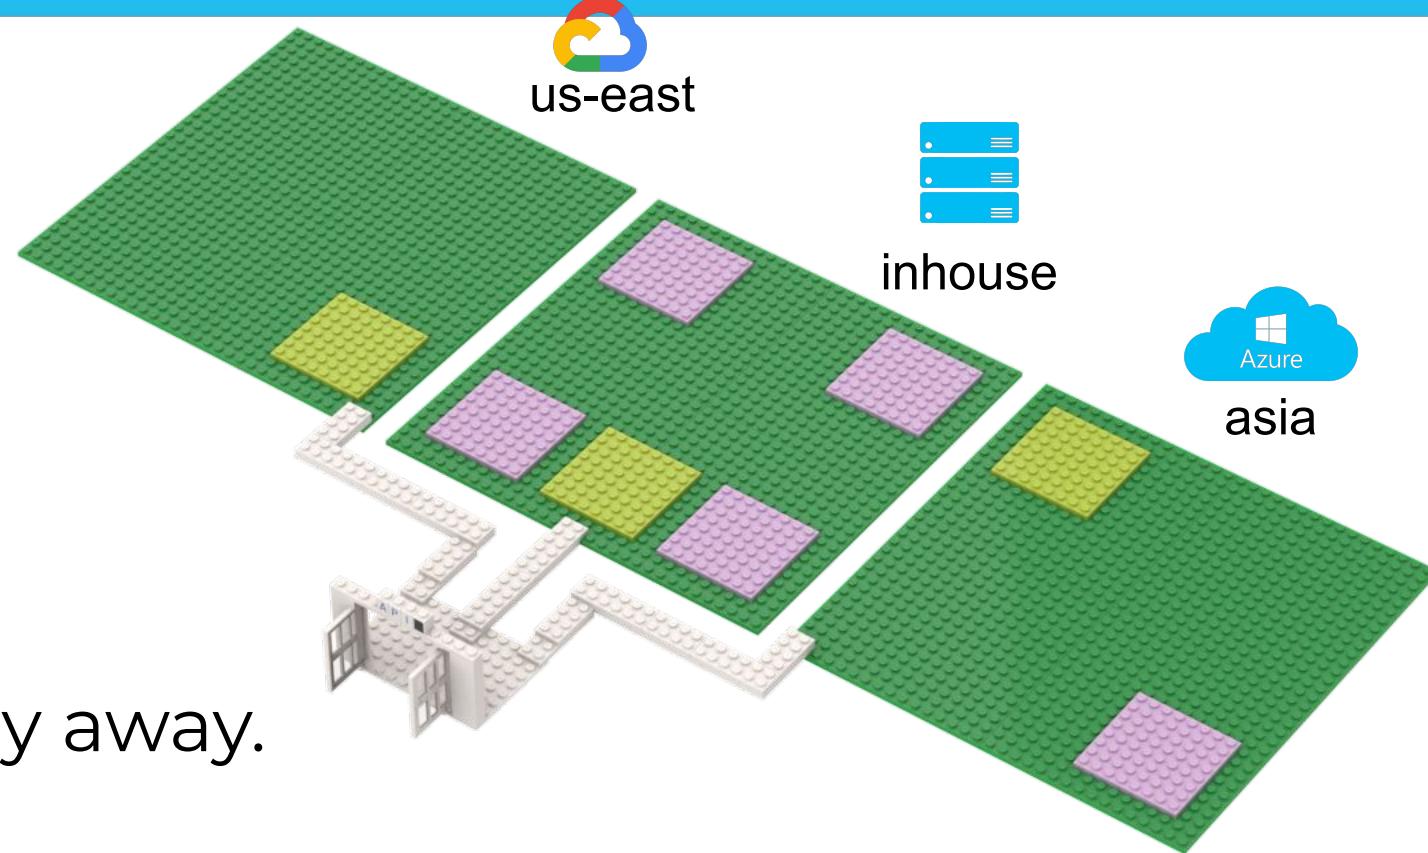
A Kubernetes for Platforms – the Platform Owner

As a platform owner

Give user a well defined,
documented platform
to ease their job!

Abstract the complexity away.

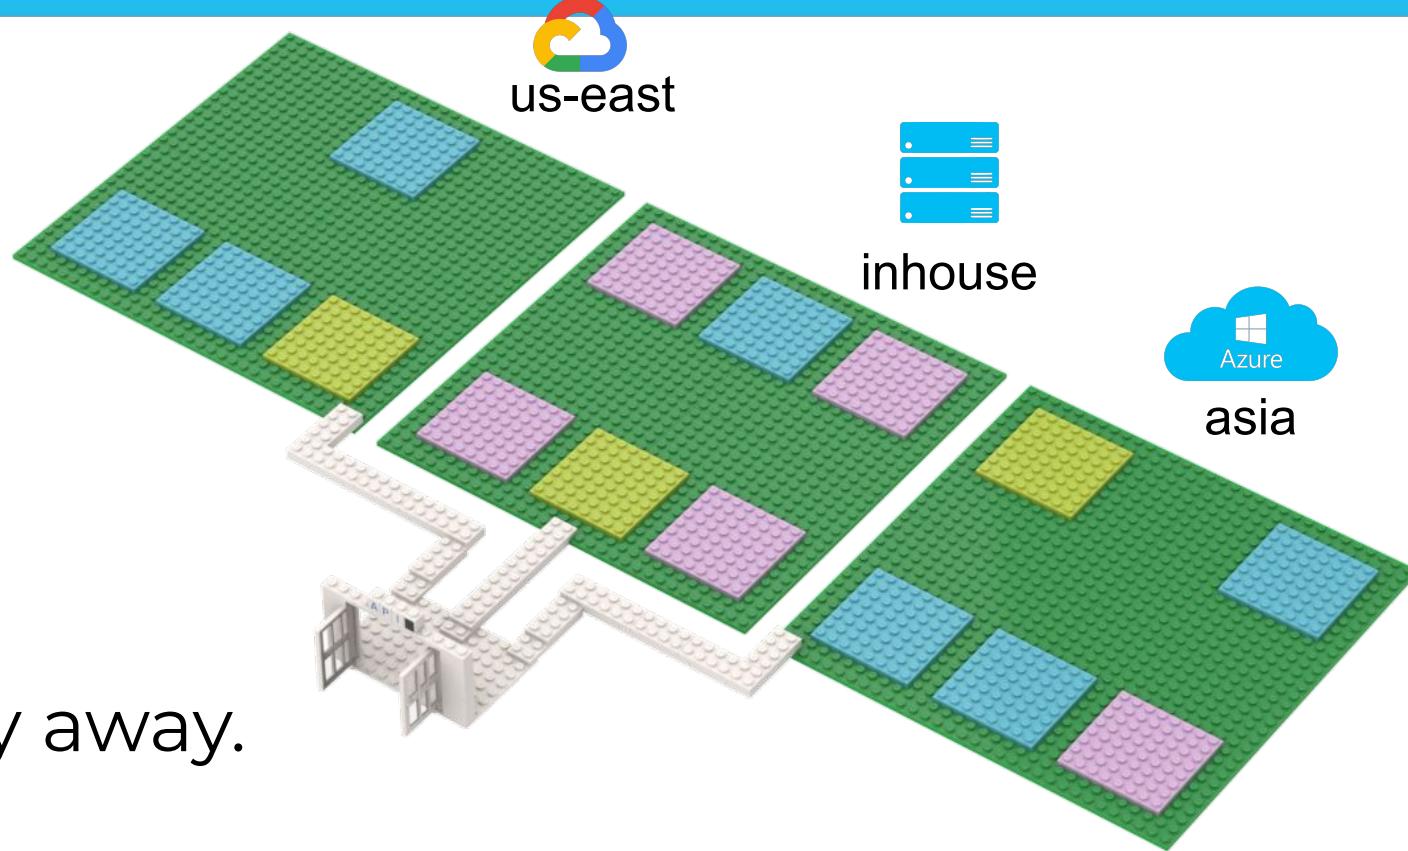
And scale it horizontally.



A Kubernetes for Platforms – the Platform Owner

As a platform owner

Give user a well defined,
documented platform
to ease their job!



Abstract the complexity away.

And scale it horizontally
without reinventing the wheel.



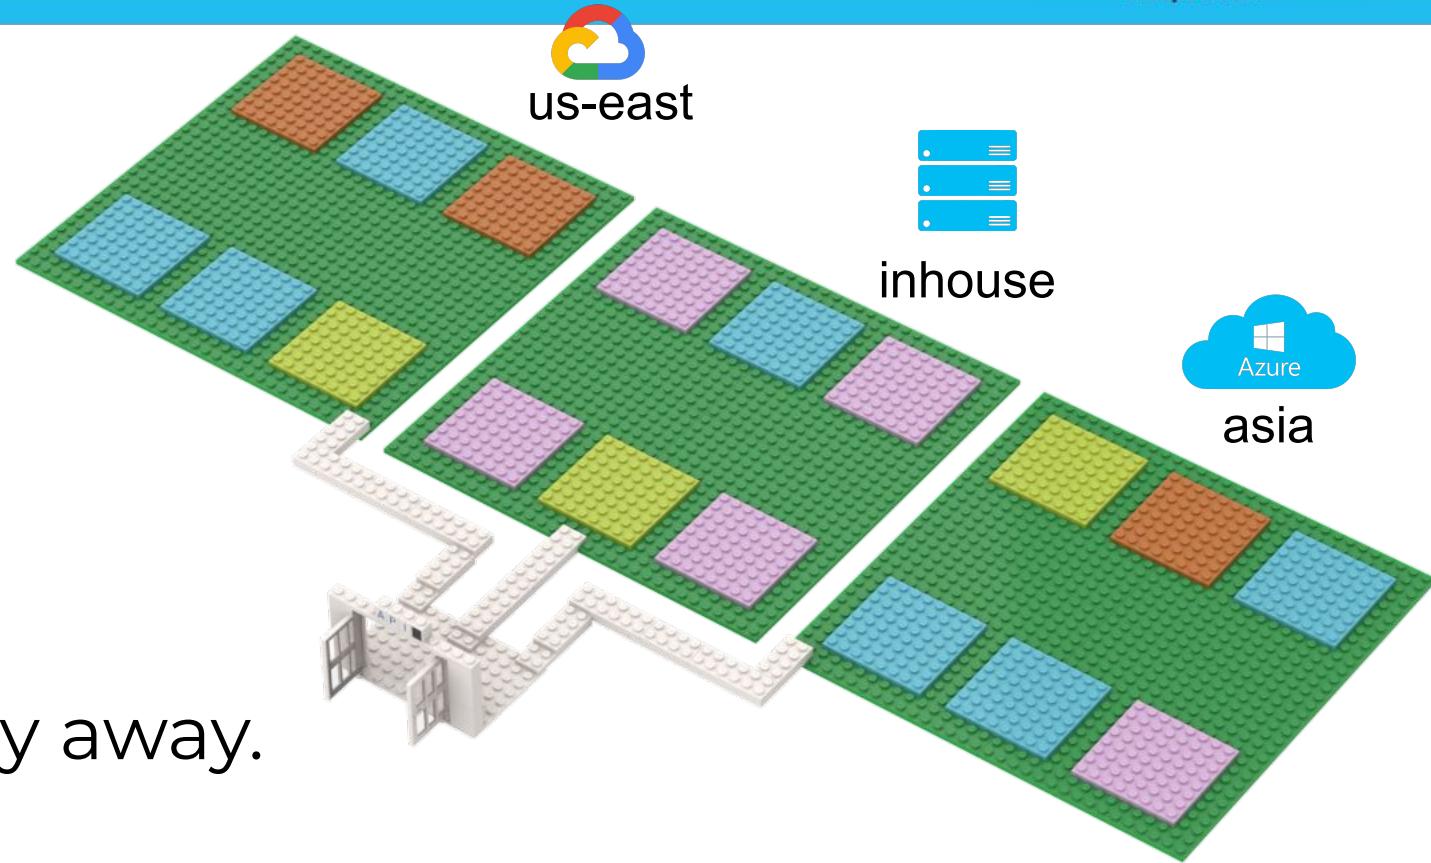
A Kubernetes for Platforms – the Platform Owner

As a platform owner

Give user a well defined,
documented platform
to ease their job!

Abstract the complexity away.

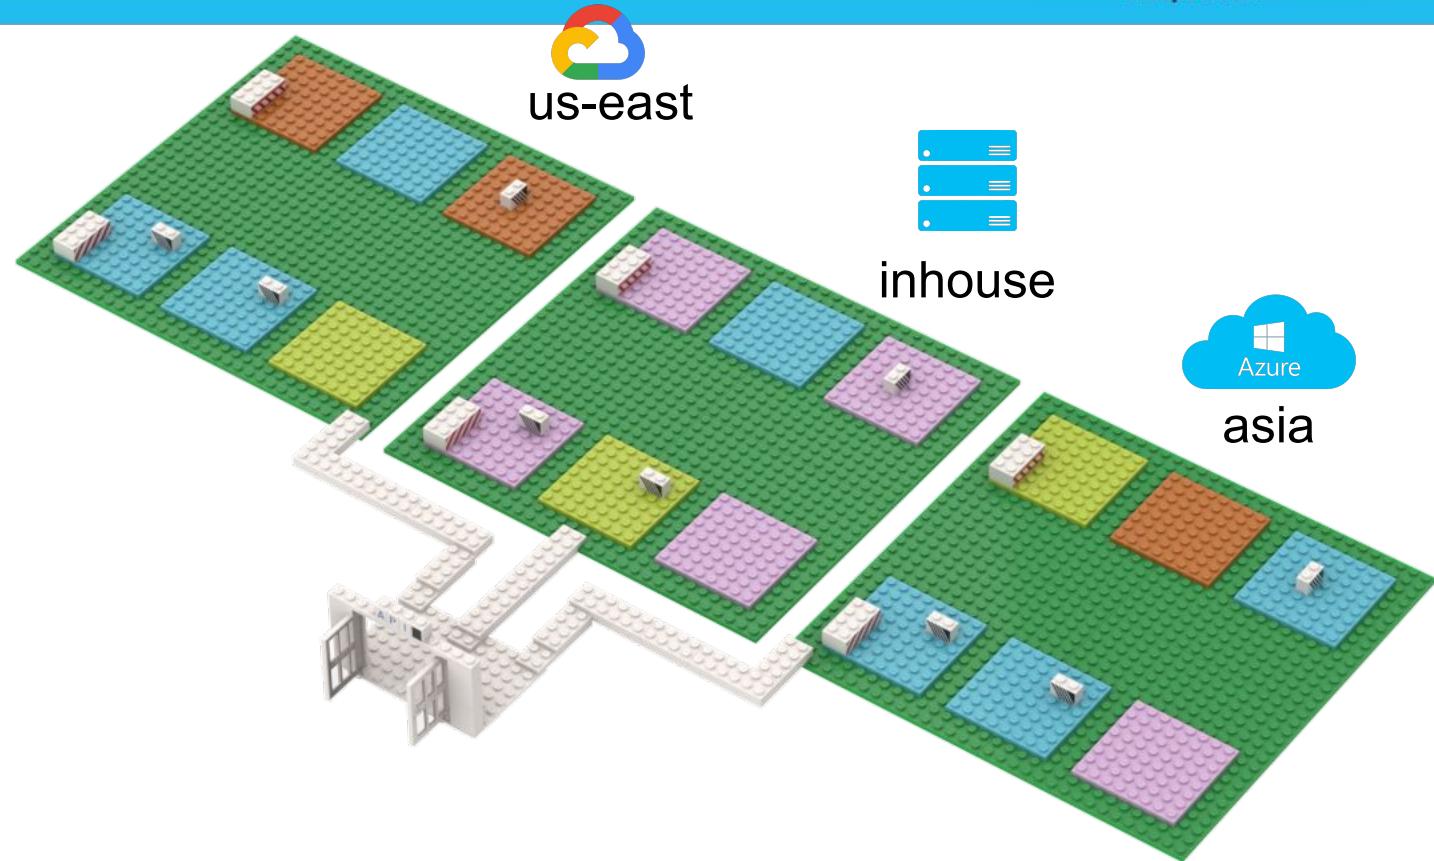
And scale it horizontally
without reinventing the wheel.



Keep it homogenous and simple.

A Kubernetes for Platforms – the Platform Owner

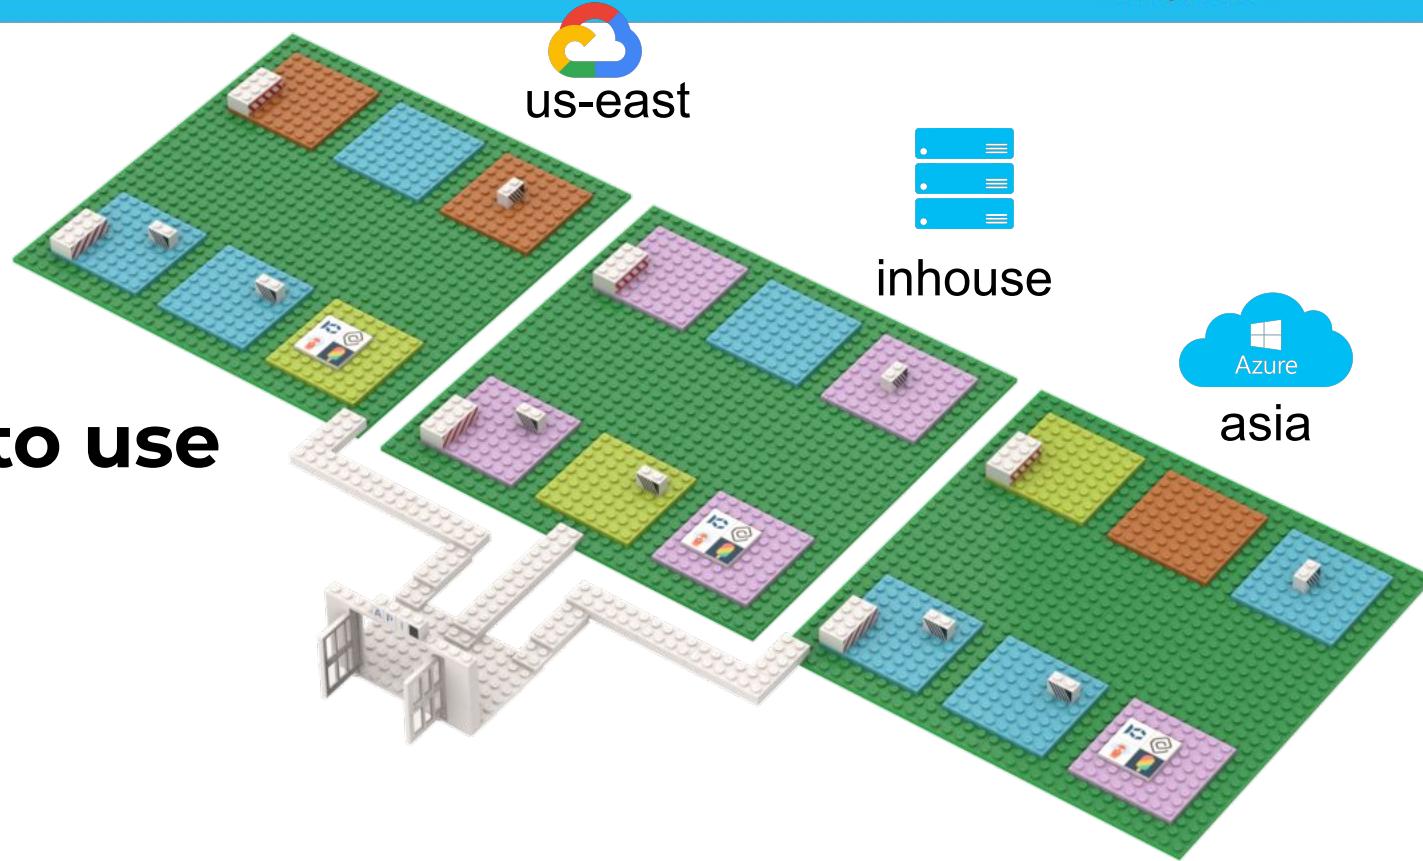
Let your users build.



A Kubernetes for Platforms – the Platform Owner

Let your users build

**You should not have
opinions on them what to use
and what not.**

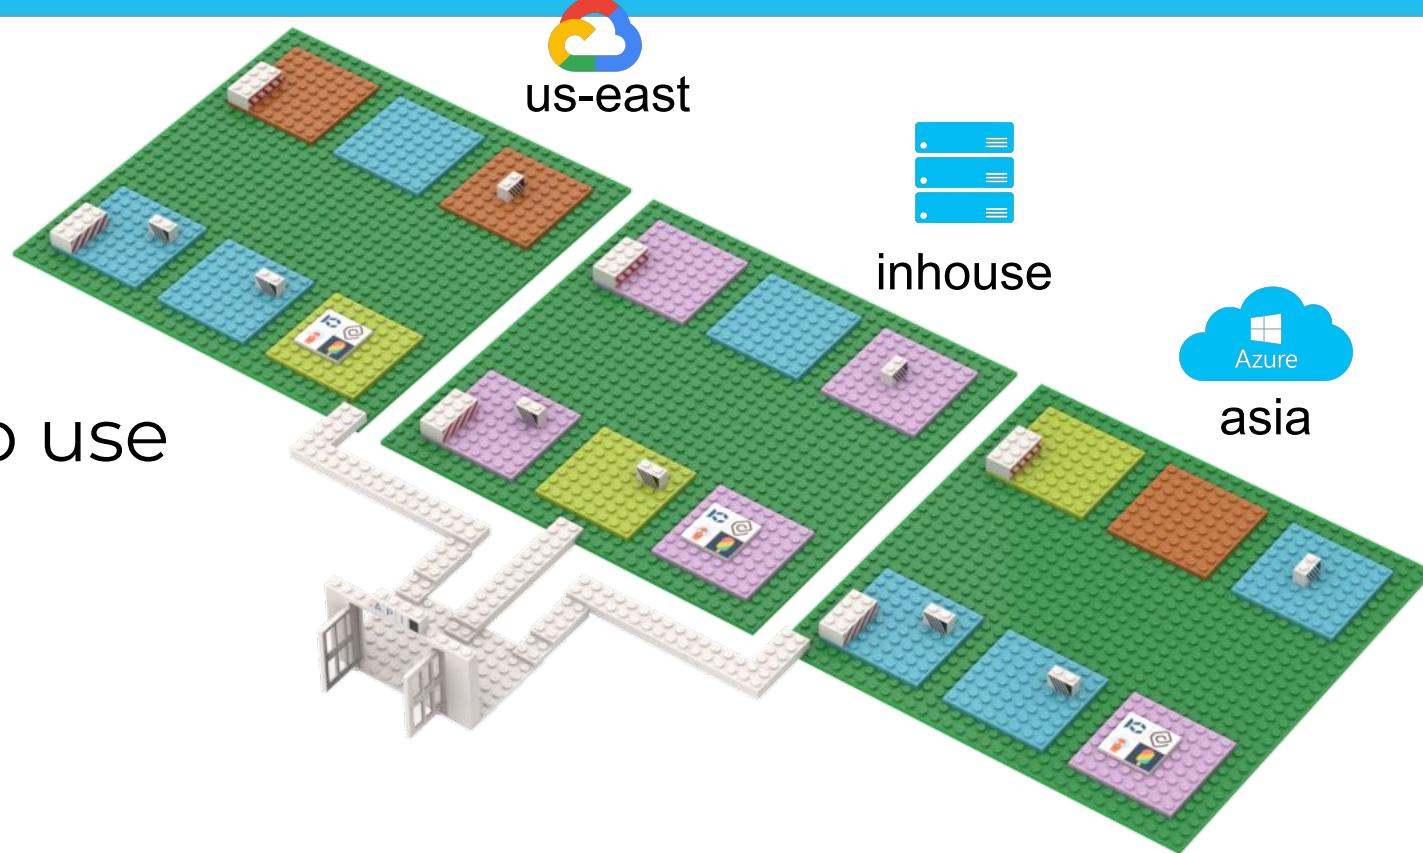


A Kubernetes for Platforms – the Platform Owner

Let your users build

You should not have
opinions on them what to use
and what not.

**Give them services
to consume.**



If they give value, your users will use them.



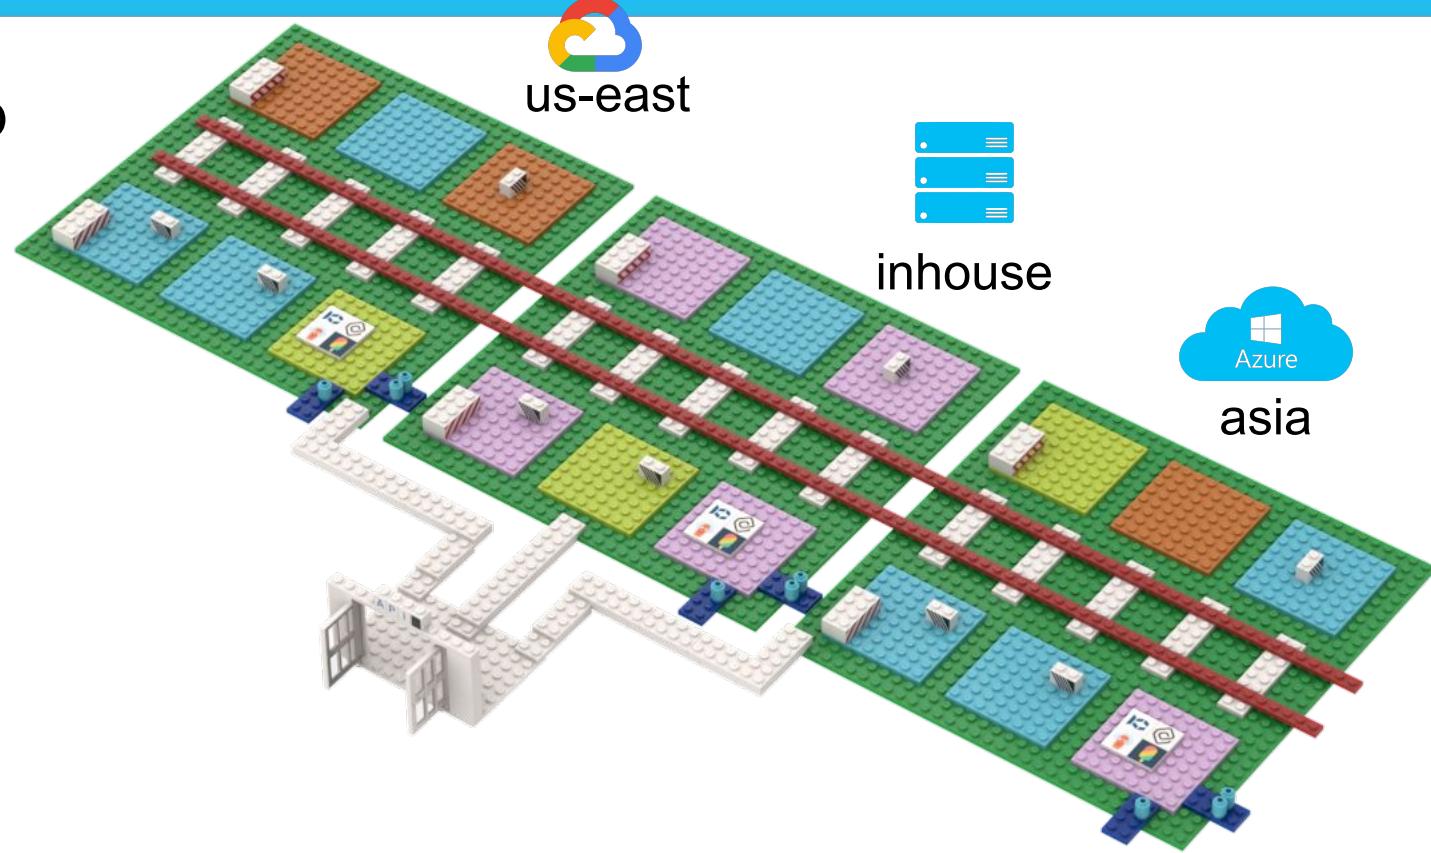


The service provider

DB team, PKI team,
CI/CD service, 3rd-parties
service owner

A Kubernetes for Platforms – the Service Provider

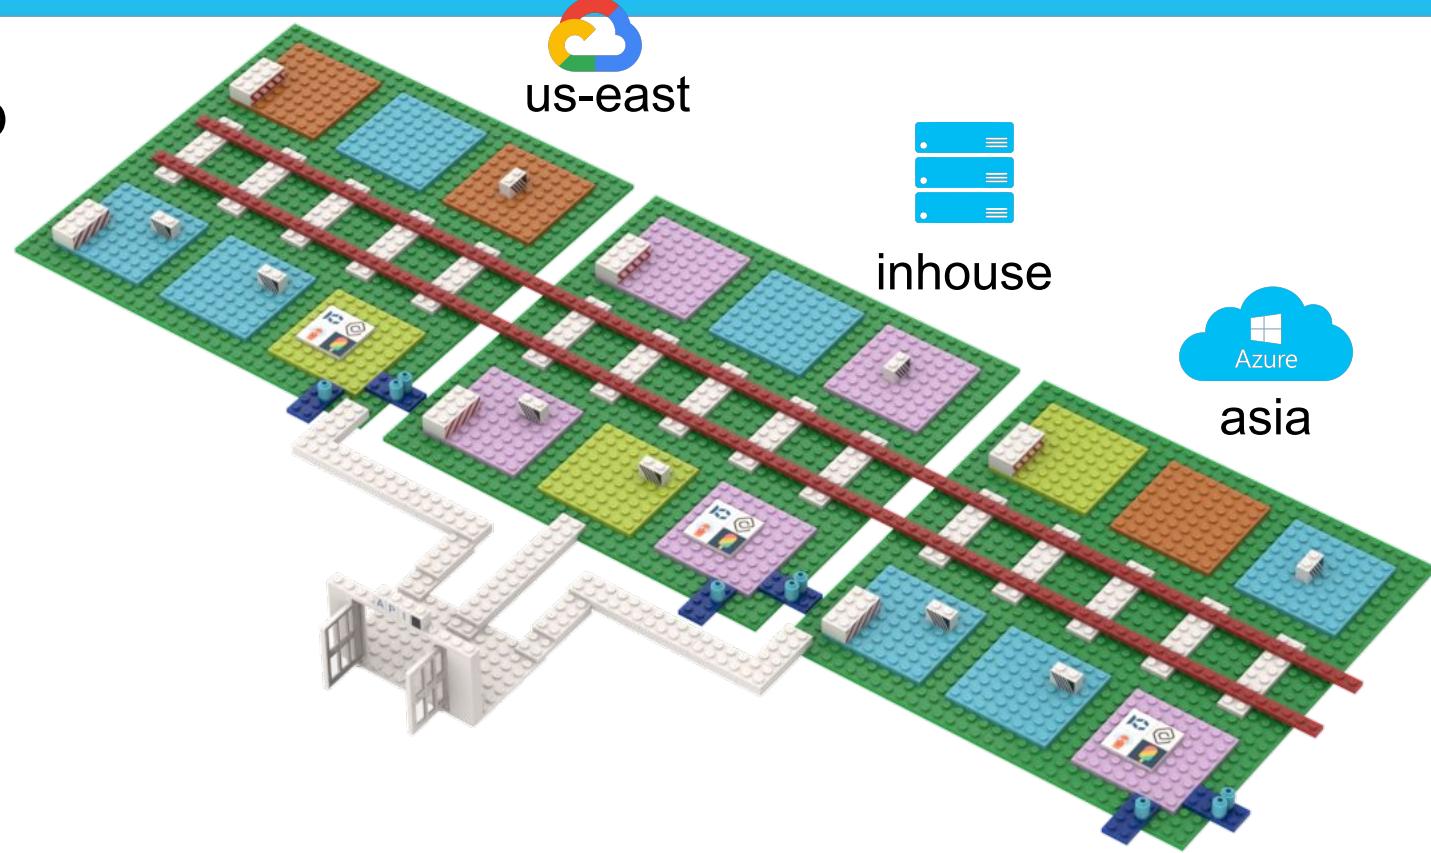
As an integrator you want to
offer a service.



A Kubernetes for Platforms – the Service Provider

As an integrator you want to offer a service.

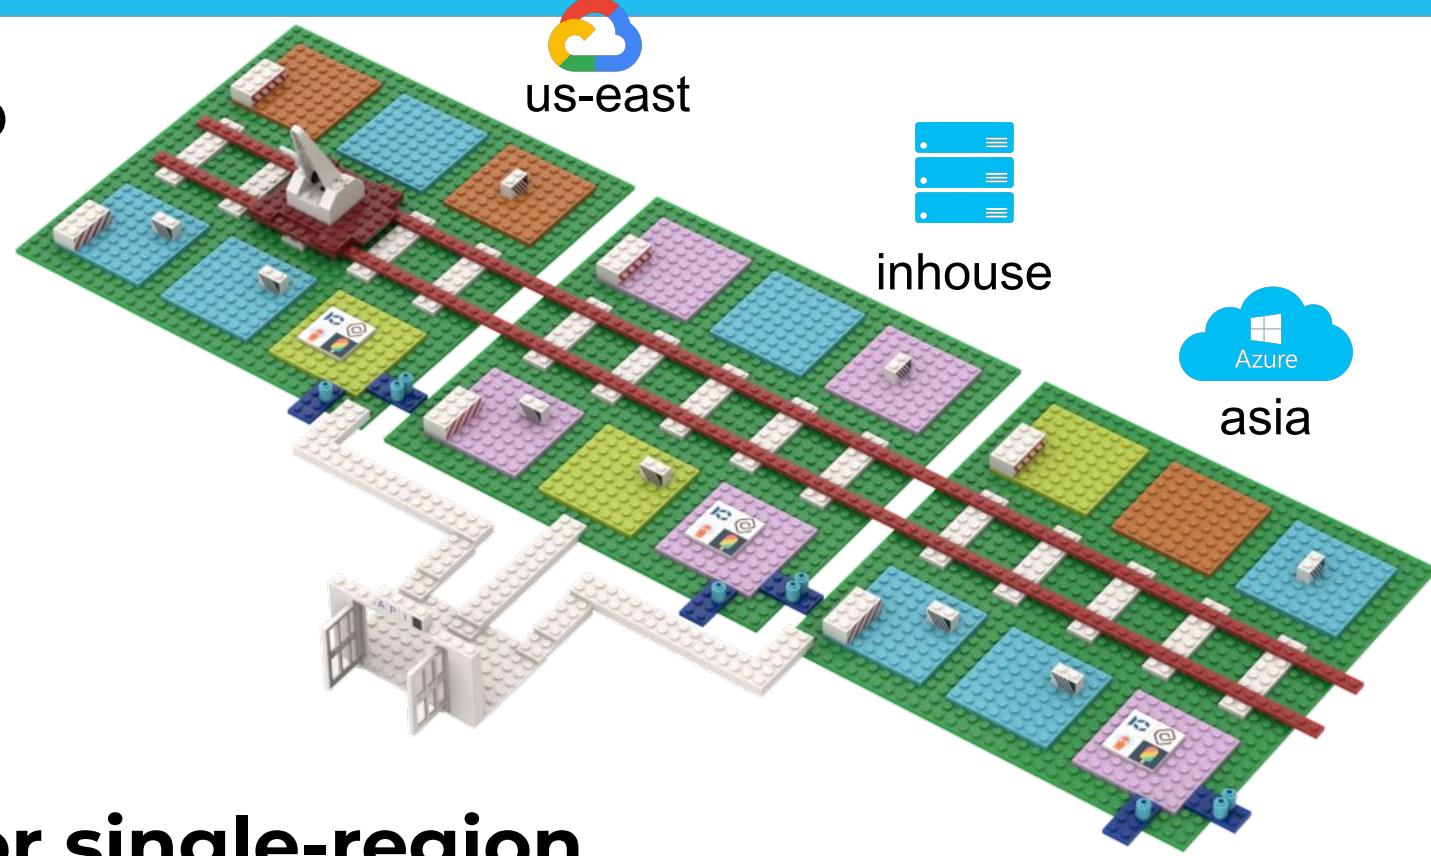
**Consistently.
Efficiently.
Safely.**



A Kubernetes for Platforms – the Service Provider

As an integrator you want to offer a service.

Consistently.
Efficiently.
Safely.



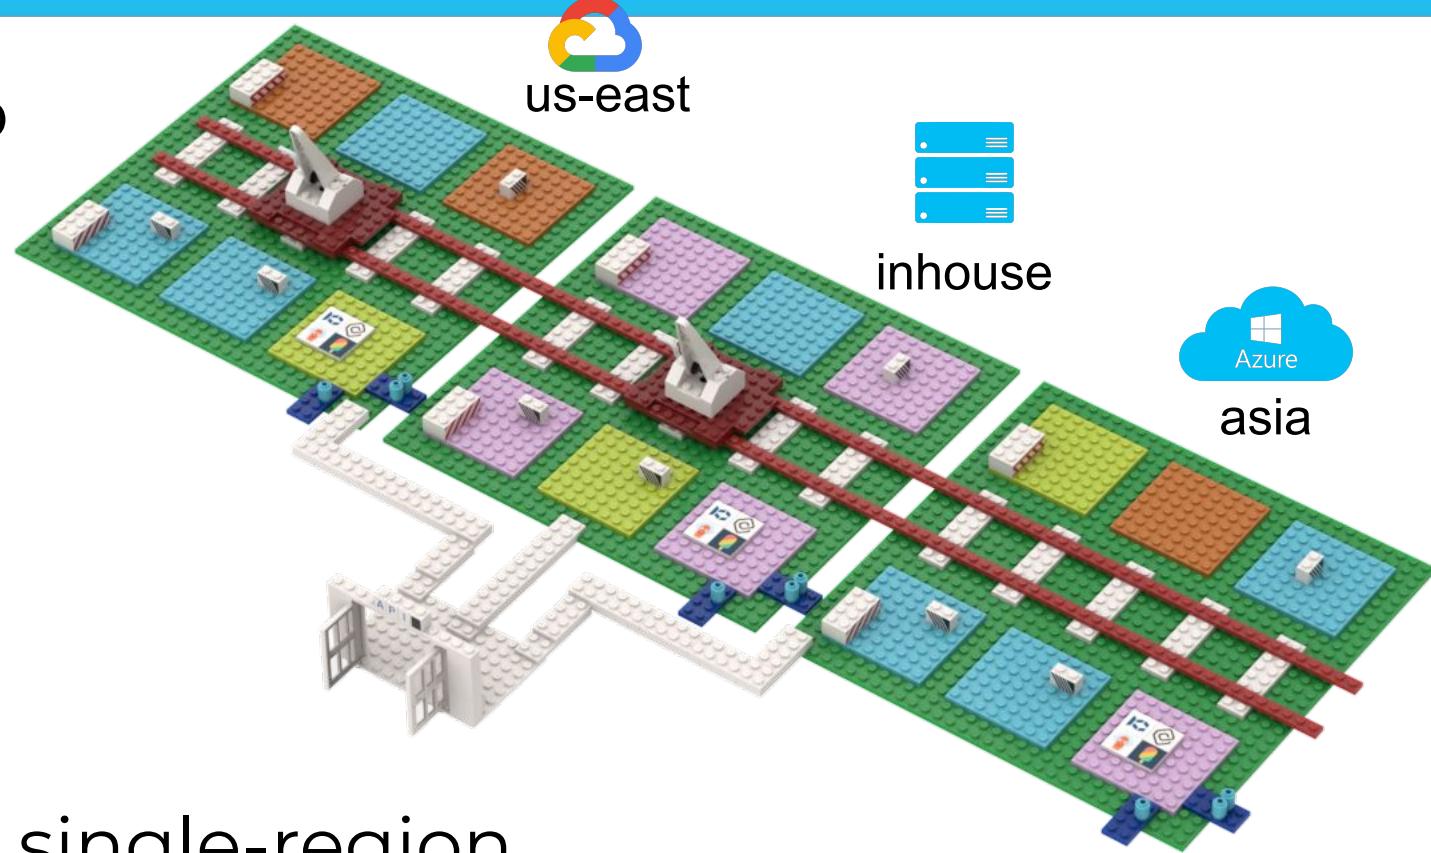
Same set of tools for single-region.



A Kubernetes for Platforms – the Service Provider

As an integrator you want to offer a service.

Consistently.
Efficiently.
Safely.



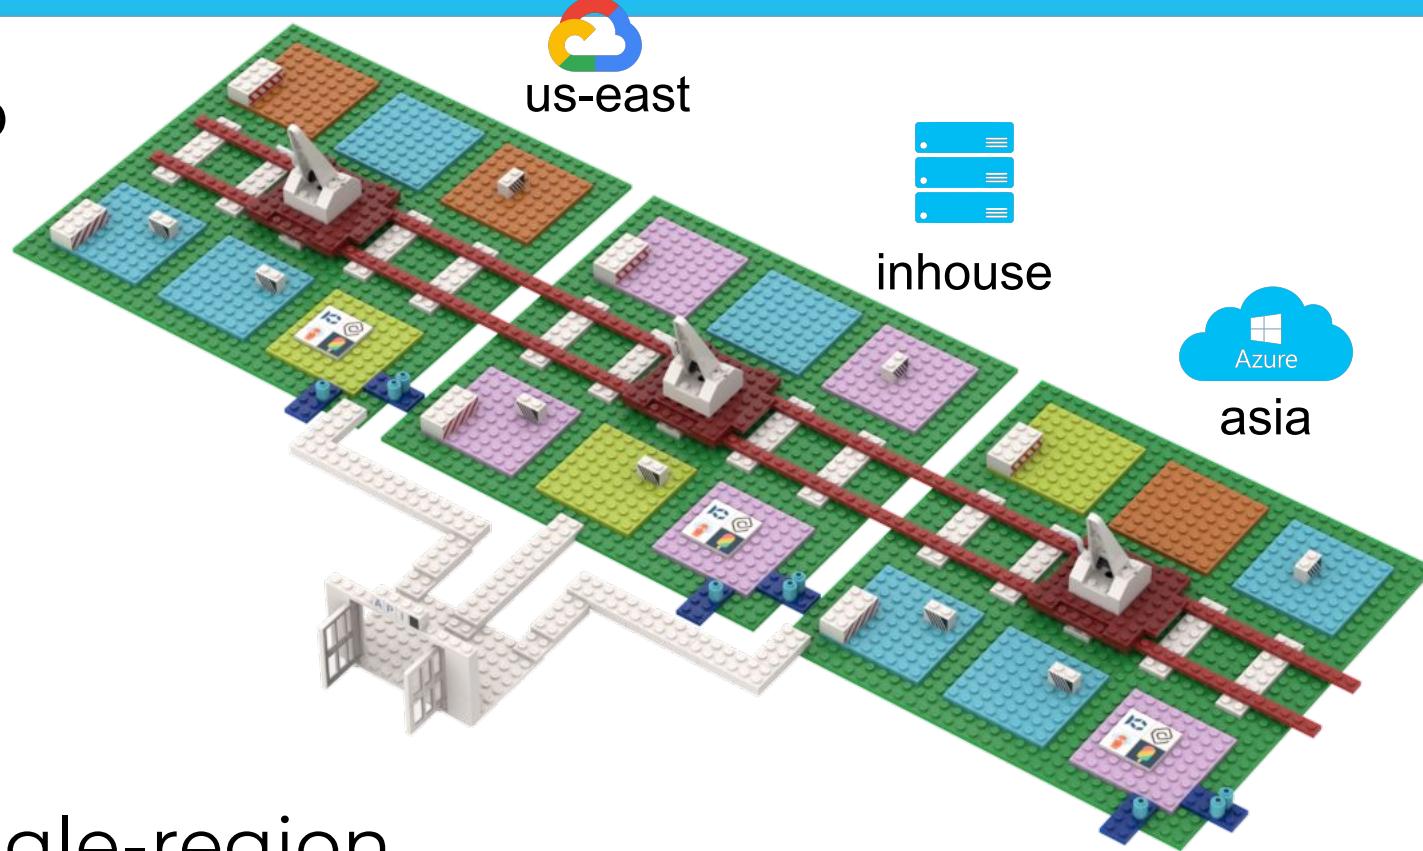
Same set of tools for single-region
or multi-region.



A Kubernetes for Platforms – the Service Provider

As an integrator you want to offer a service.

Consistently.
Efficiently.
Safely.



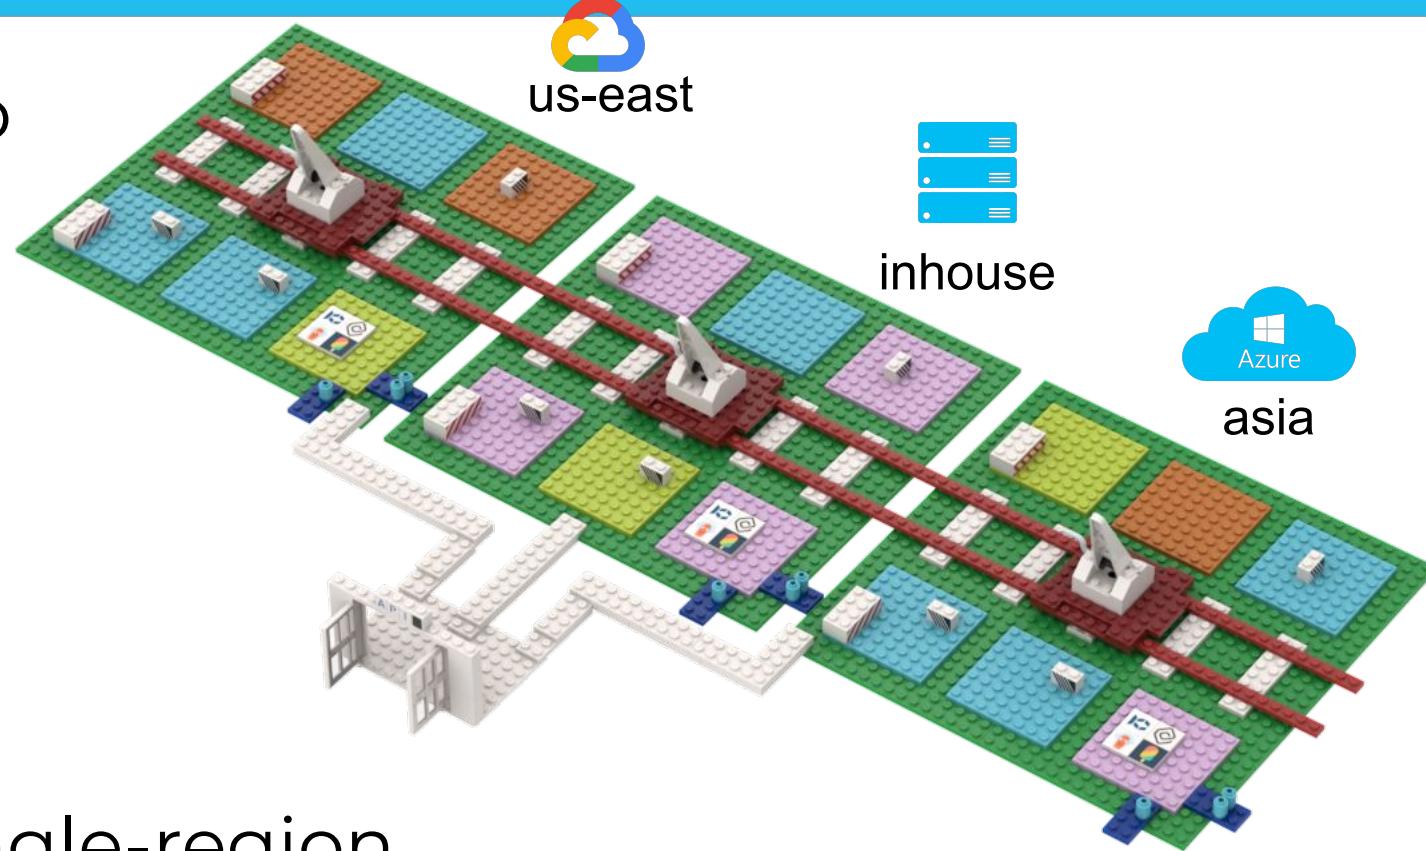
Same set of tools single-region or multi-region **or multi-cloud**.



A Kubernetes for Platforms – the Service Provider

As an integrator you want to offer a service.

Consistently.
Efficiently.
Safely.



Same set of tools single-region or multi-region or multi-cloud.

Towards a Global service as a Product.



A Kubernetes for Platforms – the Service Provider

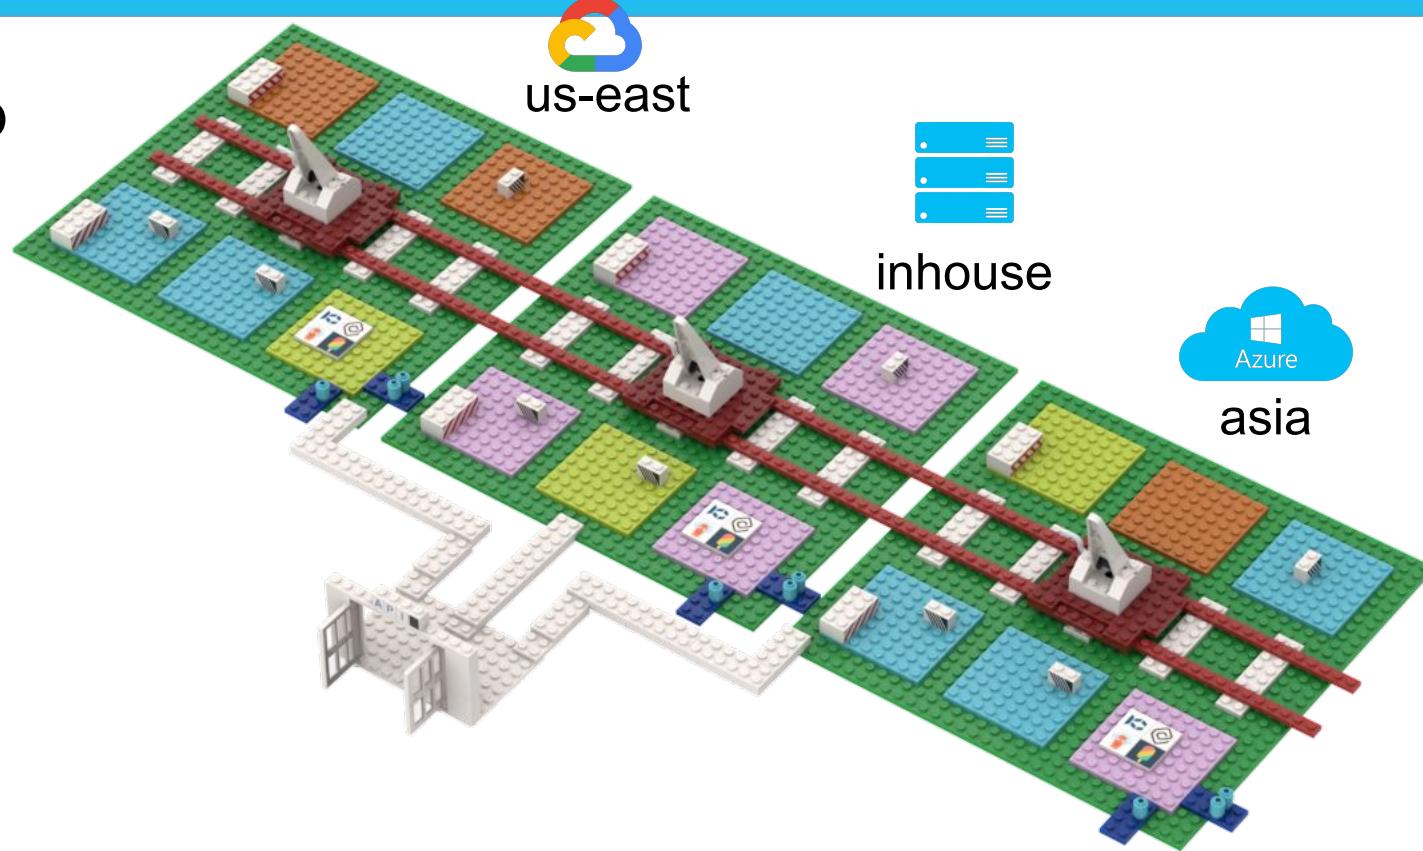
As an integrator you want to offer a service.

Consistently.

Efficiently.

Safely.

Securely.

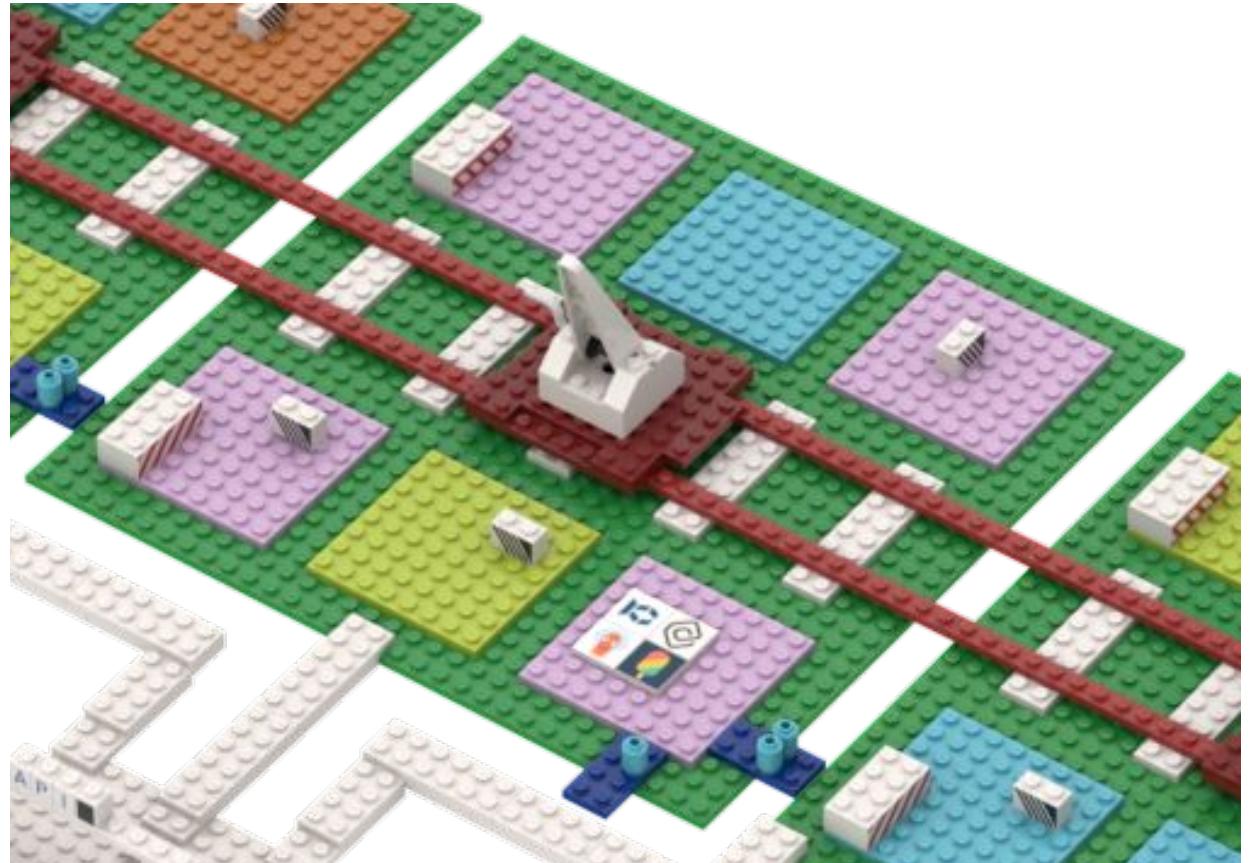


Same set of tools single-region or multi-region or multi-cloud.

Towards a Global service as a Product.



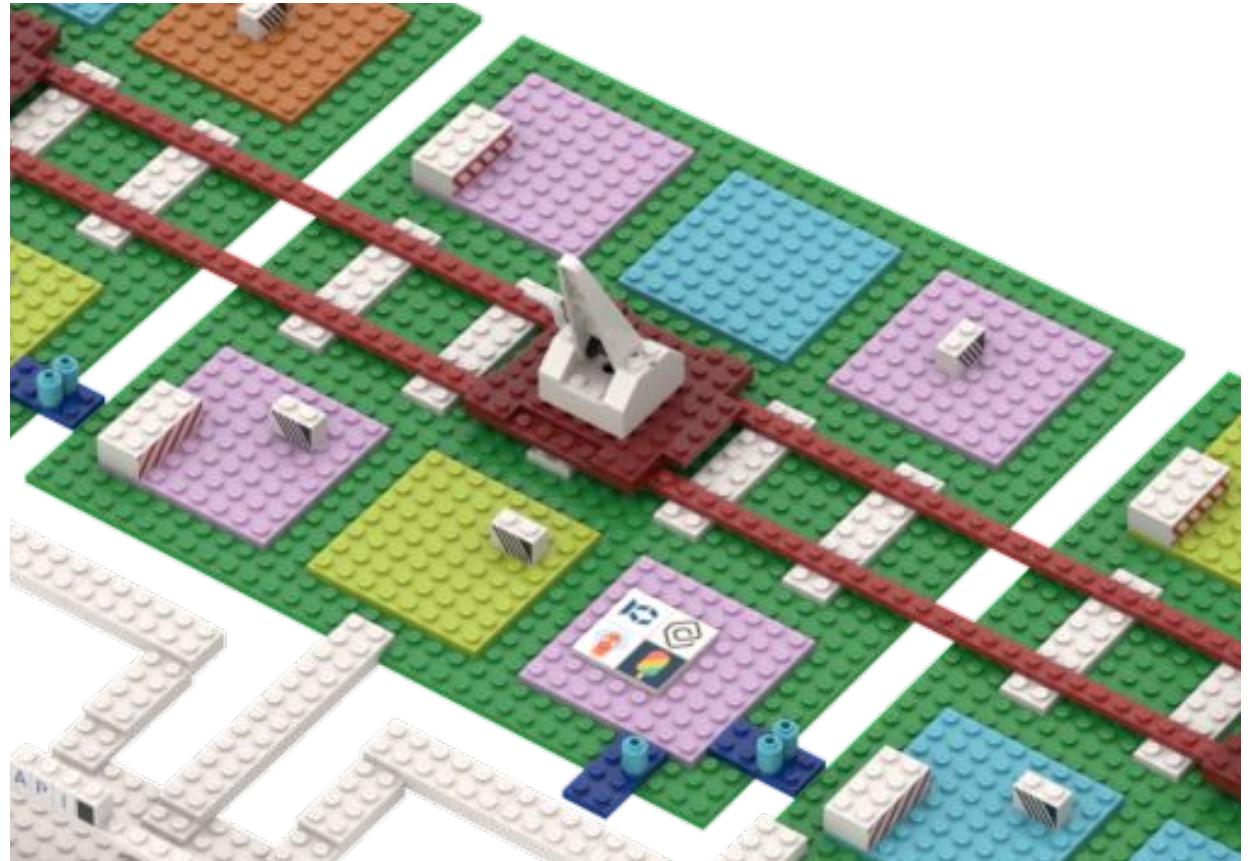
**Which primitives
does this need ?**



A Kubernetes for Platforms – the Service Provider

Which primitives
does this need ?

Are CRDs enough ?



A Kubernetes for Platforms – the Service Provider

Which primitives
does this need ?

Are CRDs enough ?

Tools like Helm ?



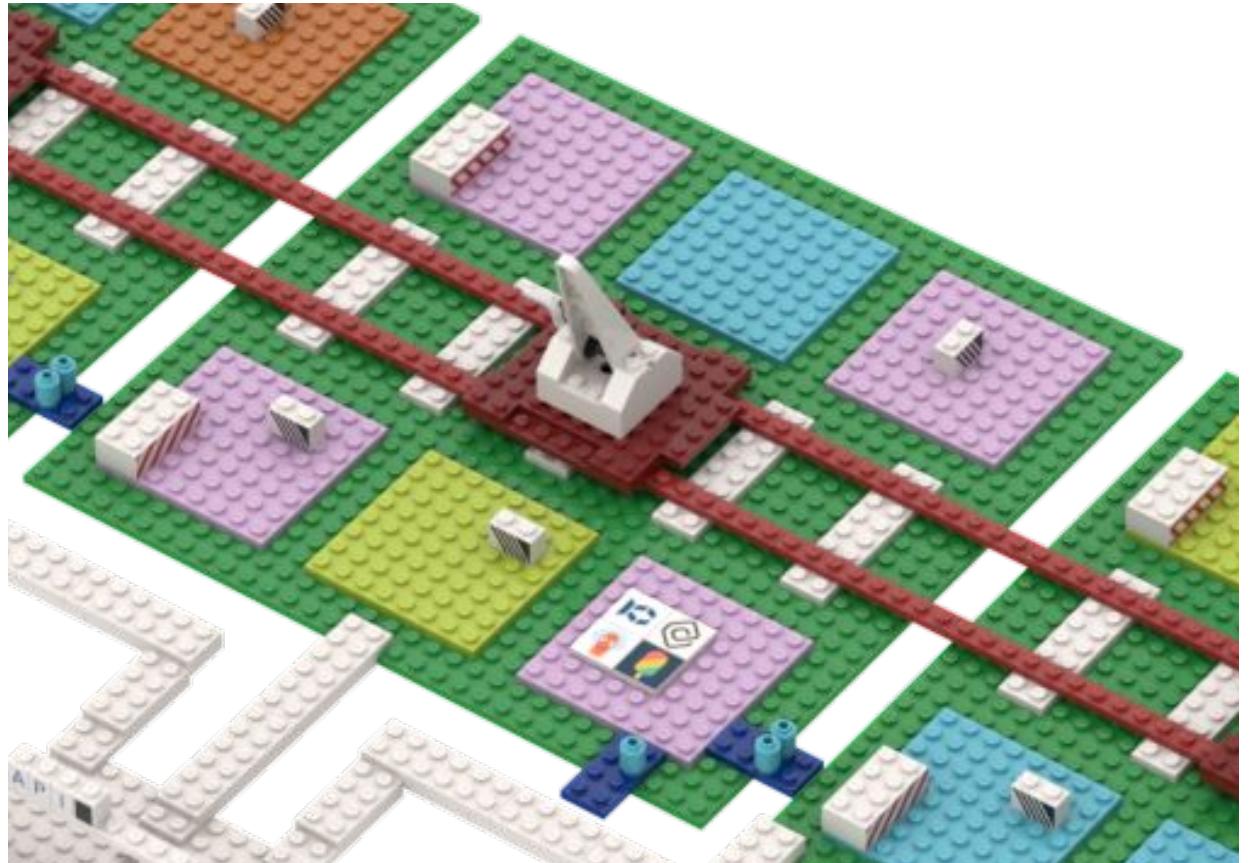
A Kubernetes for Platforms – the Service Provider

Which primitives
does this need ?

Are CRDs enough ?

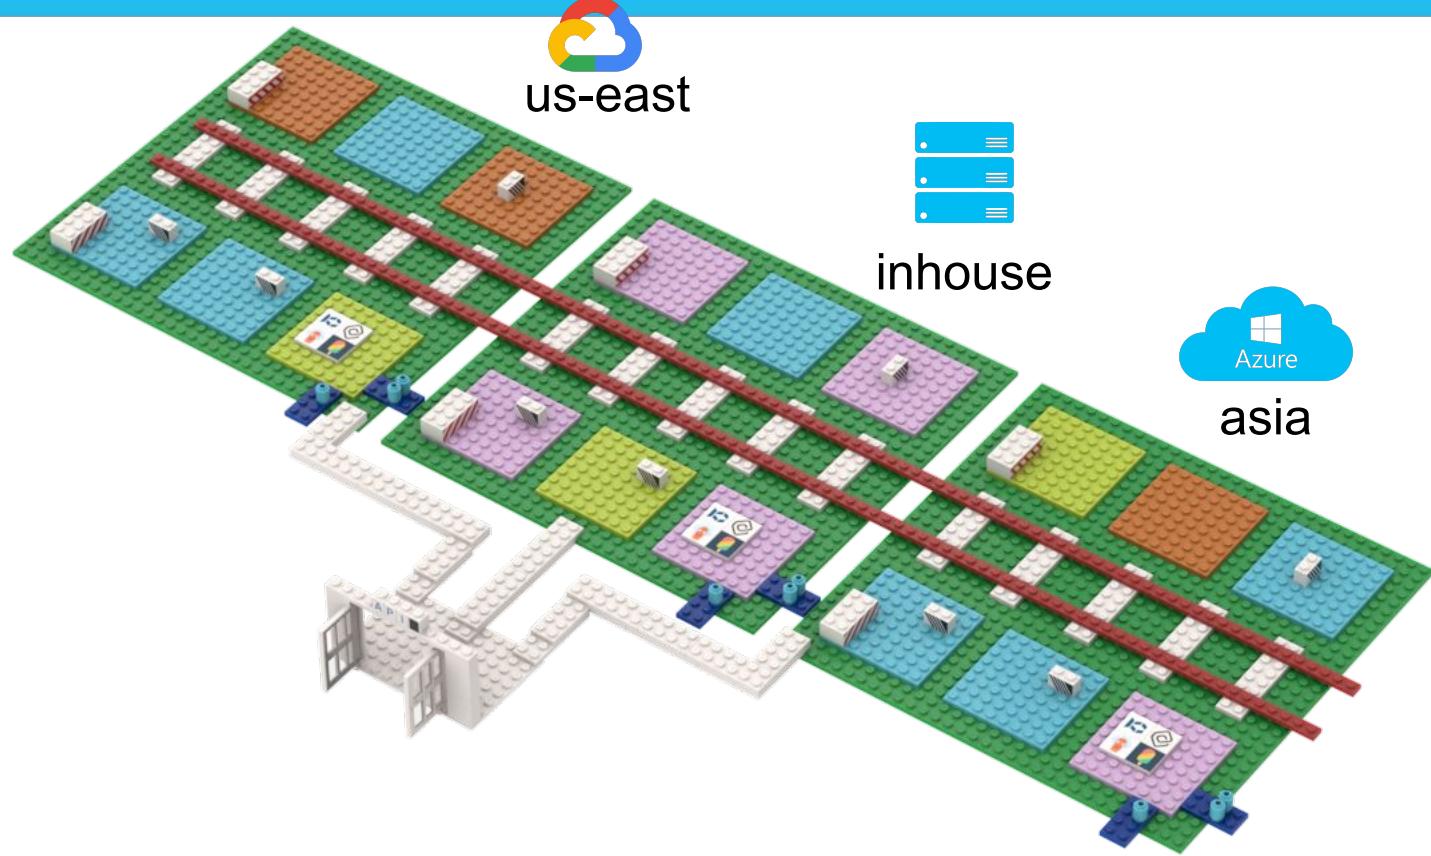
Tools like Helm ?

Kube RBAC ?



A Kubernetes for Platforms – the Service Provider

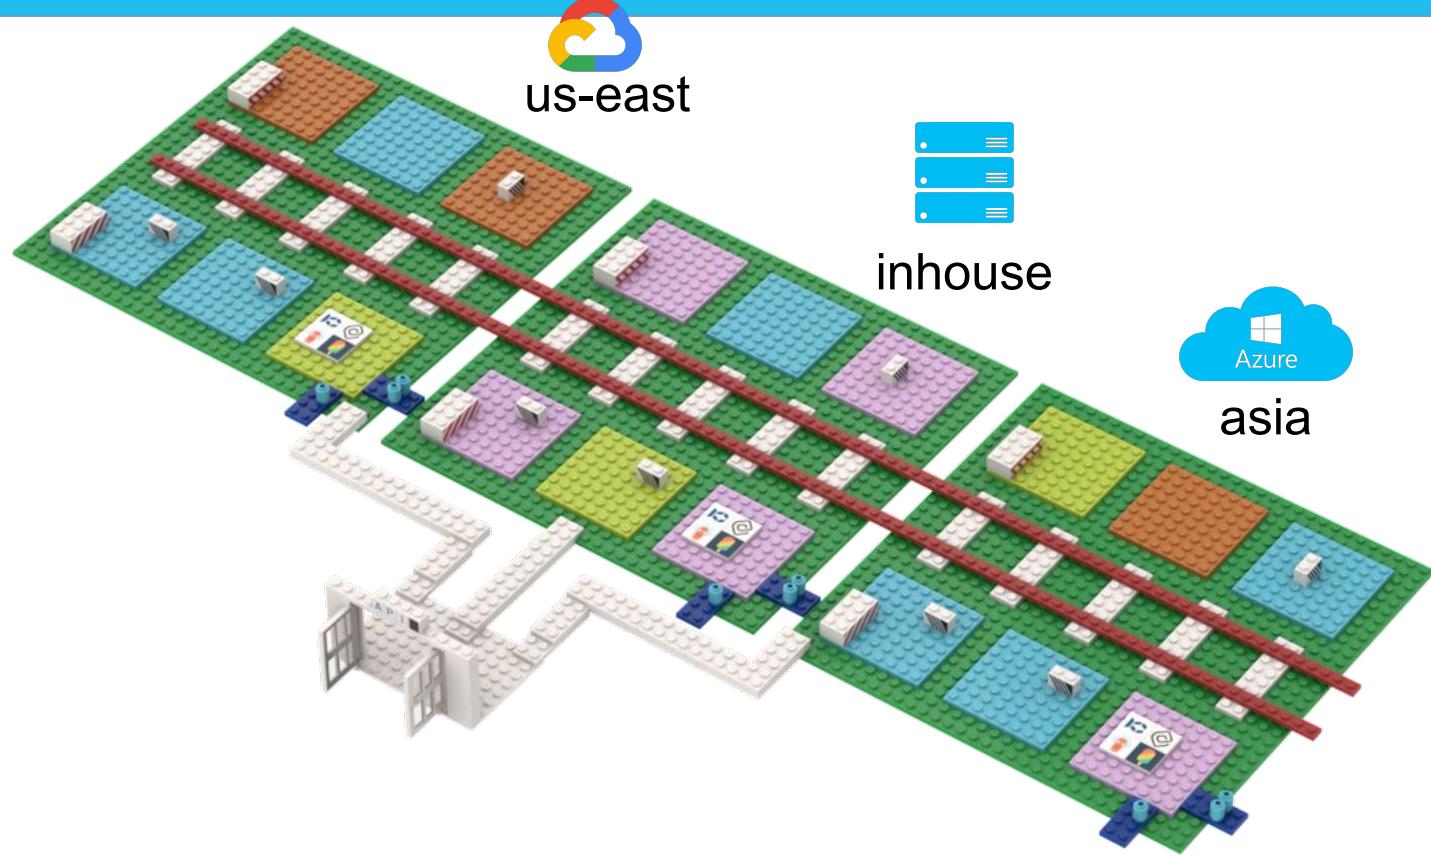
But what about compute?



A Kubernetes for Platforms – the Service Provider

But what about compute?

**Compute is just
another API.**



But what about compute?

Compute is just another API.
“Utility Computing”.



But what about compute?

Compute is just another API.
“Utility Computing”.

**We know that already with
ArgoCD or Flux or Crossplane.**



But what about compute?

Compute is just another API.
Utility Computing.

We know that already with
ArgoCD or Flux or Crossplane.

Many flavors, different APIs:
vclusters, KubeVirt, Karmada, Cluster API



But what about compute?

Compute is just another API.
Utility Computing.

We know that already with
ArgoCD or Flux or Crossplane.

Many flavors, different APIs:
vclusters, KubeVirt, Karmada, Cluster API



Compute is attached. Compute is not the center.



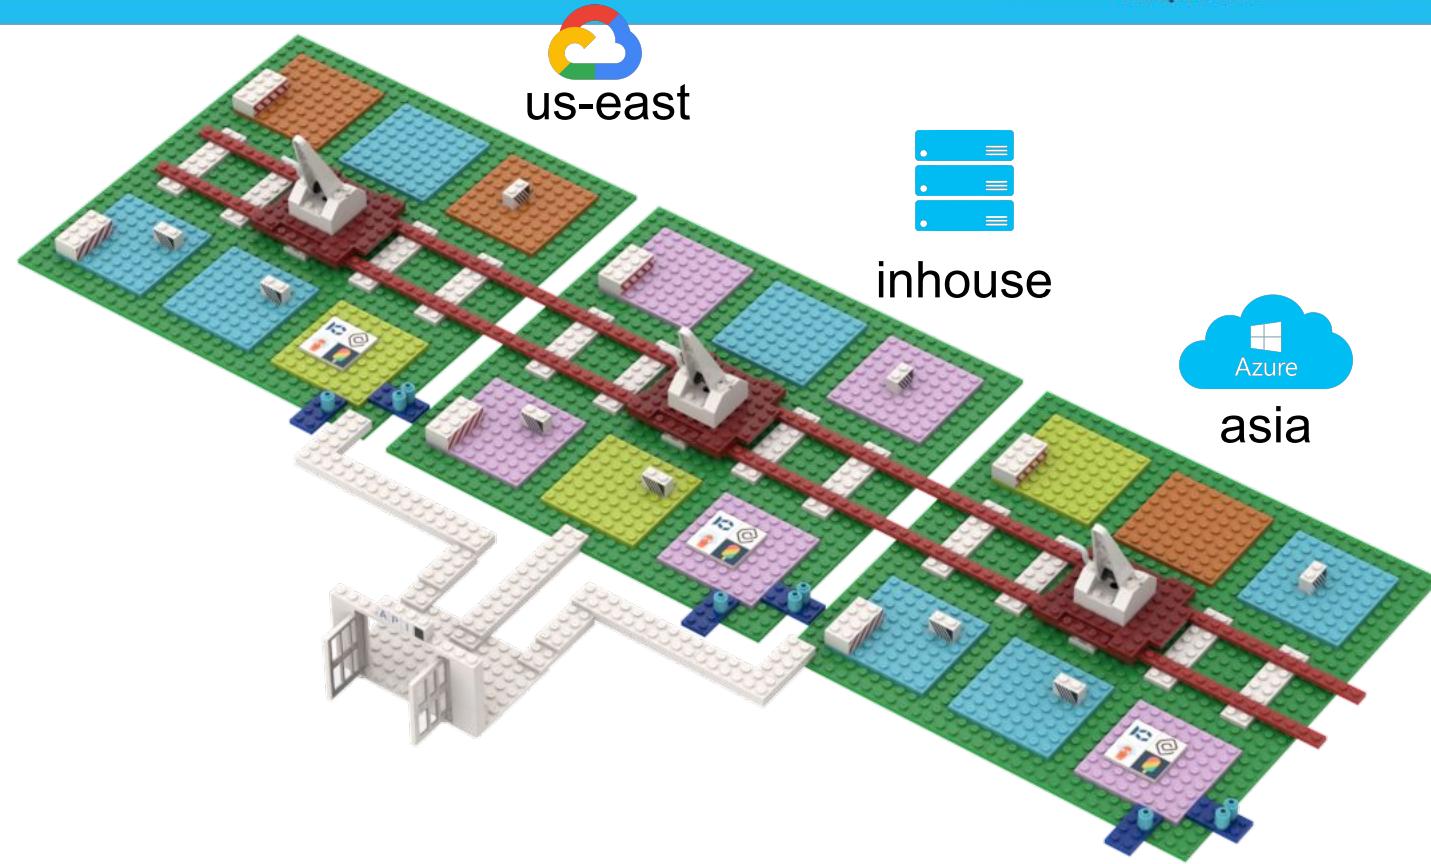


The users
Developer, application
owner

A Kubernetes for Platforms – the User

Decoupled, multi-tenant,
distributed control planes

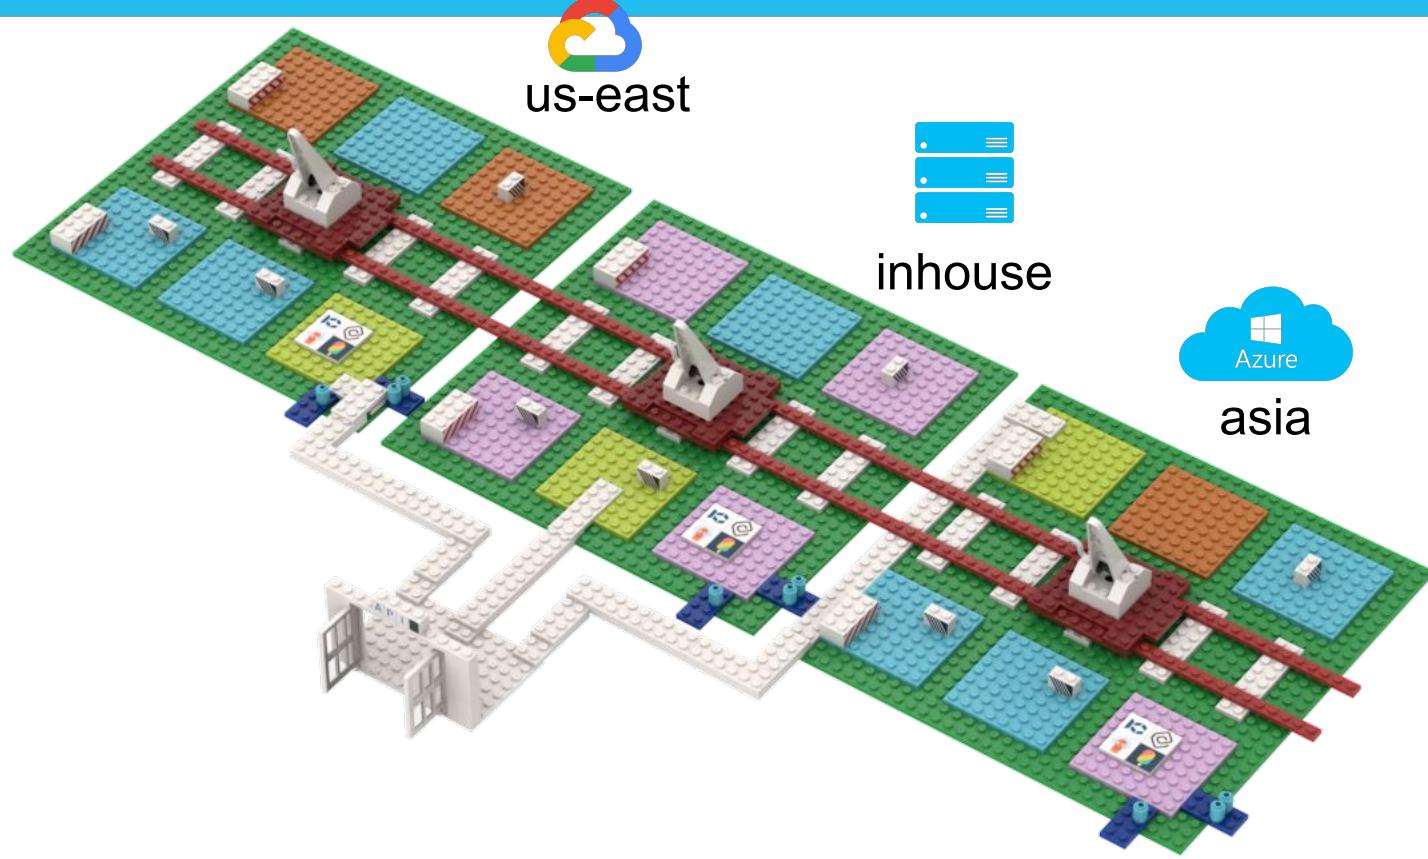
“Workspaces”



A Kubernetes for Platforms – the User

Decoupled, multi-tenant,
distributed control planes
“Workspaces”

all logically connected.

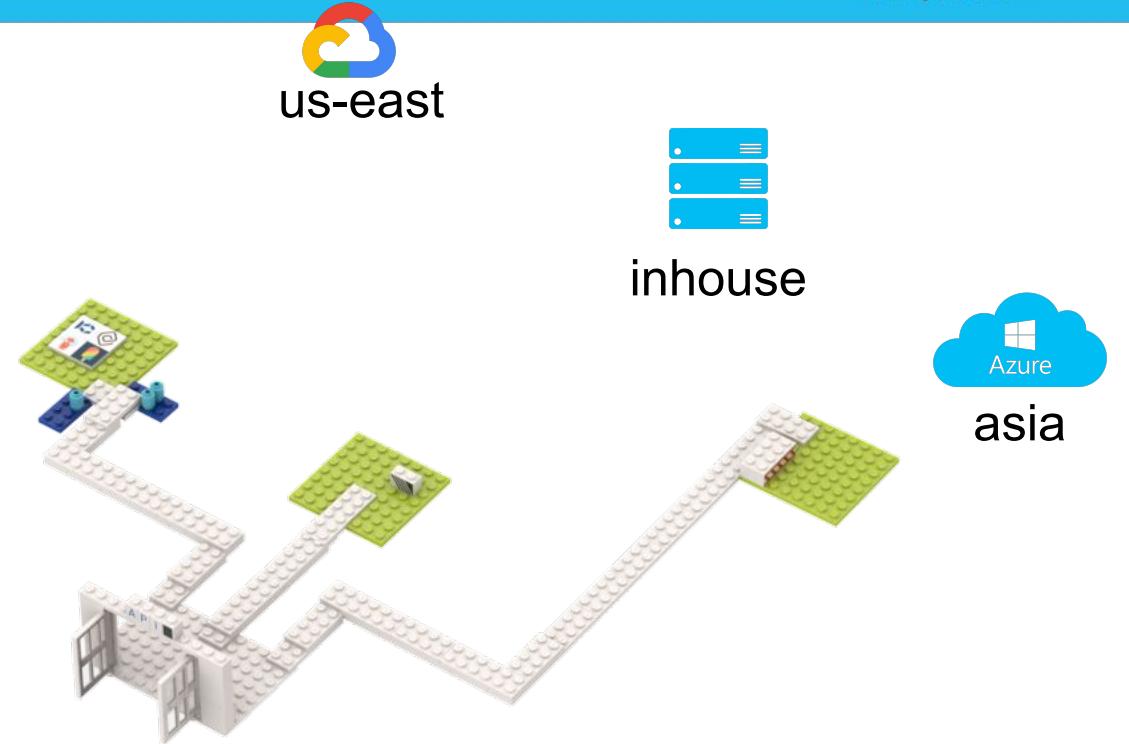


A Kubernetes for Platforms – the User

Decoupled, multi-tenant,
distributed control planes
“Workspaces”

all logically connected.

Easy to navigate.



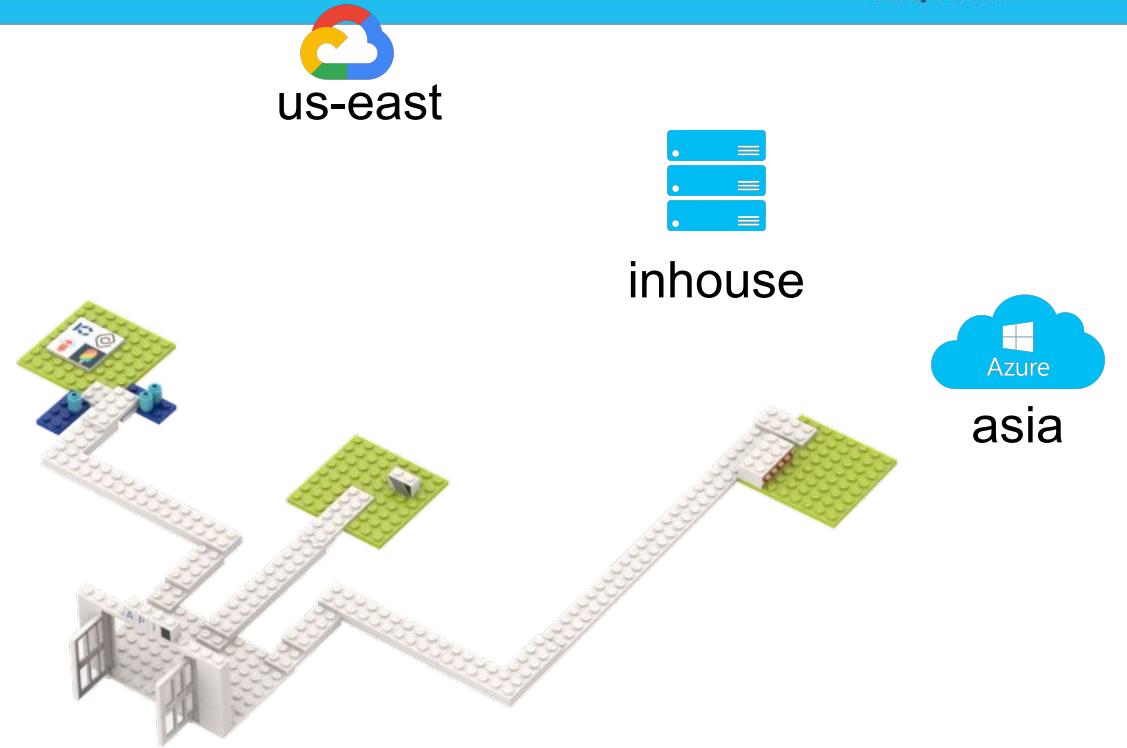
A Kubernetes for Platforms – the User

Decoupled, multi-tenant,
distributed control planes
“Workspaces”

all logically connected.

Easy to navigate.

from app to app
from region to region
from cloud to cloud.



What we described is not Kubernetes.
You cannot build it just with more clusters.
We need new primitives.
New primitives for a world that is not just containers
for a world that is inherently multi-tenant/region/cloud.
New primitives to regain speaking the same language.
New primitives to collaborate on together.

KCP - Kubernetes-like Control Planes

- CNCF Sandbox project
- Framework to bootstrap Kubernetes-like SaaS/Platforms
- Work in progress to move “generic controlplanes” to upstream



Demo

```
# region:name:root = Europe
```

```
# region:name:beta = US
```

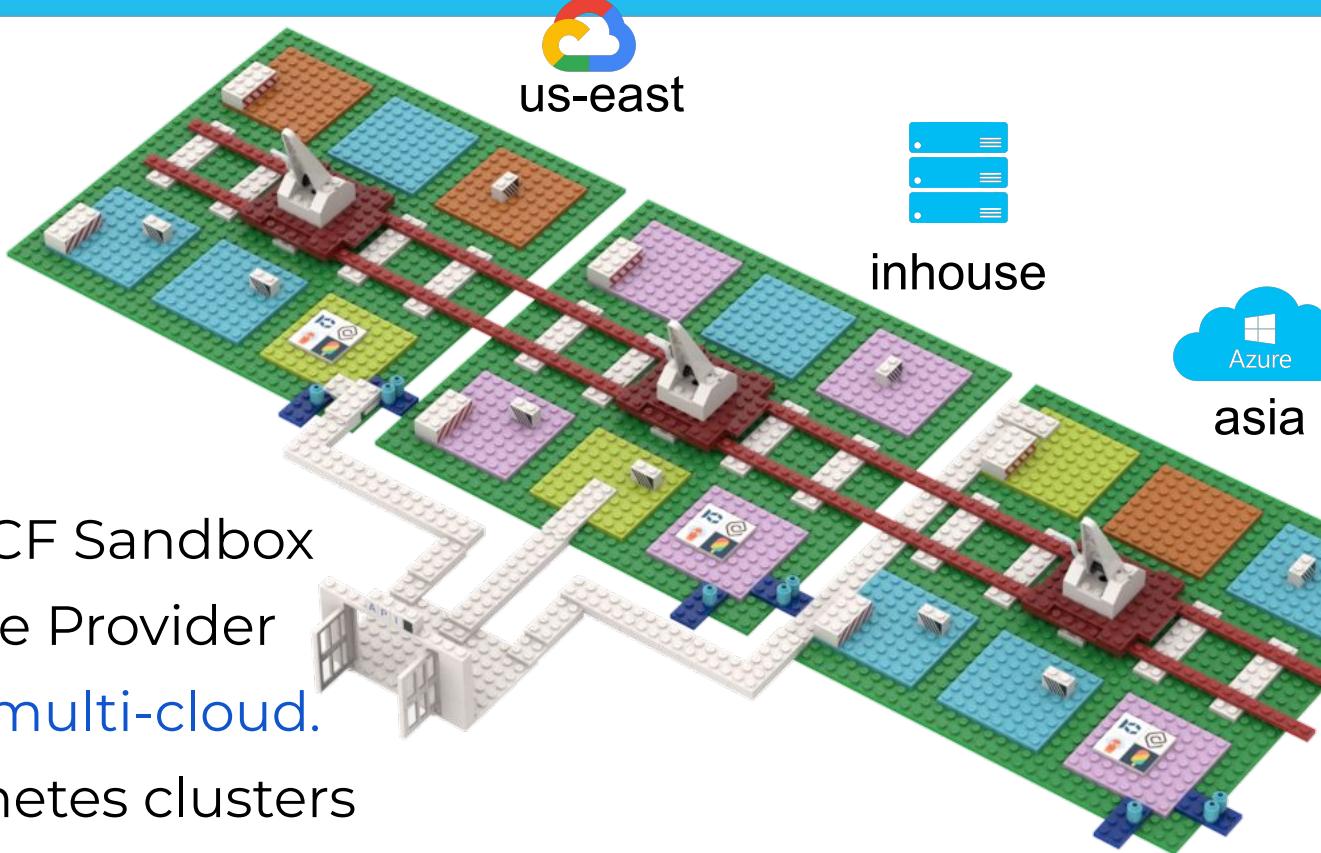
```
$ |
```

KCP – a framework exploring these ideas



kcp

- Kubernetes like APIs
- Based on Kubernetes source code, CNCF Sandbox
- Platform Personas: Owner, User, Service Provider
- Inherently multi-tenant, multi-region, multi-cloud.
- Workspaces, each like (generic) Kubernetes clusters
- All under one API endpoint.
- Horizontally scalable (sharding)
- Extended API management for service providers





Talk to us - look for logo! We have stickers!



Find us at #kcp-dev kubernetes slack



github.com/kcp-dev/kcp



@kcp @the_sttts @sscheele @mangirdas



G14



L11



E7

Find us at one of these booths!



Slack channel QR

Feedback



Thanks